

INSURANCE VERSUS MORAL HAZARD IN INCOME-CONTINGENT STUDENT LOAN REPAYMENT*

Tim de Silva[†]
MIT Sloan

[Click Here for Latest Version](#)

September 26, 2023

Abstract

This paper studies the trade-off between providing insurance and disincentivizing labor supply in student loans with income-contingent repayment. Using discontinuities in repayment rates from Australia's student loan system, I show borrowers adjust their labor supply to reduce repayments on income-contingent loans. These responses are larger in occupations with more hourly flexibility, among young borrowers with more debt, and among liquidity-constrained borrowers with less wealth and larger housing payments. I use these responses to estimate a structural model and find they are consistent with a Frisch labor supply elasticity of 0.11 and substantial frictions that limit labor supply adjustment. In this model, a constrained-optimal income-contingent loan generates welfare gains relative to a fixed repayment contract equivalent to a 1.3% increase in lifetime consumption, with the same fiscal cost. Equity contracts generate gains that are larger on average but significantly more dispersed. The labor supply responses to income-contingent repayment reduce the insurance these contracts can provide at a given cost, but they are too small to justify fixed repayment contracts.

***Incomplete draft: Please do not circulate.** I'm grateful to my dissertation committee, Taha Choukhmane, Jonathan A. Parker (co-chair), Lawrence D.W. Schmidt, Eric C. So, and David Thesmar (co-chair) for their continuous support. I thank Pat Adams, David Autor, Maxime Bonelli, Bruce Chapman, Hui Chen, Marc de la Barrera, Maryam Farboodi, Amy Finkelstein, Brice Green, Jonathan Gruber, Sebastian Hillenbrand, Tong Liu, Debbie Lucas, Pierfrancesco Mei, Lira Mota, Christopher Palmer, Jim Poterba, Charlie Rafkin, Antoinette Schoar, Andrei Shleifer, Kerry Siani, Adam Solomon, Yevhenii Usenko, Emil Verner, Adrien Verdelhan, Rodrigo Verdi, Constantine Yannelis, and seminar participants at the Inter-Finance PhD Seminar, MIT Sloan, and MIT Economics for helpful comments and discussions. I also thank the ATO *ALife* and ABS DataLab teams for extensive data assistance, in particular Andrew Carter, Justin Holland, and Son Nguyen, Simone Melchionna for sharing his Fortran expertise, Andrew Norton for supplying data and responding to my never-ending stream of questions, and Jenn Mace for her unwavering support. Access to the data required for this project would not have been possible without the help of Anna Bedford, Nicholas Biddle, Andrew Norton, and Joe Weber. This research was supported by the Alfred P. Sloan Foundation through a Household Finance grant awarded to the National Bureau of Economic Research, MIT Sloan School of Management, Mark Kritzman and Elizabeth Gorman Research Fund, Stone Finance Ph.D Fund, Thomas Anthony Pappas Endowed Scholarship, MIT SuperCloud, and the Lincoln Laboratory Supercomputing Center. I'm grateful to the ANU Centre for Social Research and Methods and University of Technology Sydney for hosting me as a visiting scholar. The results of these studies are based, in part, on tax data supplied by the Australian Taxation Office to the Australian Bureau of Statistics under the Taxation Administration Act 1953. The final page of this paper contains a full required disclaimer. All remaining errors are my own.

[†]Contact information: www.timdesilva.me, tdesilva@mit.edu.

In many countries, students finance higher education through government-provided student loans. These loans are the second-largest household liability in the US at \$1.6 trillion and account for 10% of household debt in the US and UK. Traditionally, government-provided student loans have resembled debt contracts, where borrowers make fixed payments after graduation to repay their loan balances. Because student loans are generally not dischargeable in bankruptcy, these contracts force individuals to bear most of the risk associated with the returns to higher education. Unfortunately, the risk of low income upon graduation materializes for many borrowers, with 25% of US borrowers defaulting within five years after graduation (Hanson 2022).

One potential policy to provide more insurance against income risk is to make student loans more equity-like by indexing payments to borrowers' incomes. This idea has been discussed extensively (Friedman 1955; Shiller 2004; Palacios 2004; Chapman 2006; Zingales 2012), and governments in the US, UK, Canada, and Australia have recently implemented it by providing income-contingent loans. In contrast to non-dischargeable debt contracts, income-contingent repayment provides insurance by reducing payments as a borrower's income declines. However, this insurance potentially comes at the cost of creating moral hazard: because repayments increase with income, borrowers have an incentive to reduce labor supply to decrease repayments. Empirically, income-contingent repayment appears effective at providing insurance (Mueller and Yannelis 2019; Herbst 2023), but there is no consensus on the moral hazard effects it creates (Yannelis and Tracey 2022).


The objective of this paper is to study two central questions. First, how and through what mechanisms does income-contingent repayment affect borrowers' labor supply? Second, what form of income-contingent repayment in a government-provided financing contract optimally balances the cost of moral hazard with the benefits of providing insurance? To identify labor supply responses empirically, I leverage administrative data and policy variation from the Australian Higher Education Loan Programme (HELP), the first program to provide income-contingent loans nationwide. I then use these responses to estimate a structural life cycle model and study the implications of various forms of income-contingent repayment. In my normative analysis, I consider a social planner that maximizes borrower welfare, taking education and borrowing choices as given.

My main empirical finding is that borrowers reduce their labor supply to lower repayments on income-contingent loans. These responses are larger in occupations with more hourly flexibility, among young borrowers with more debt, and among liquidity-constrained borrowers. However, my structural estimation shows these responses are quantitatively small: replicating this evidence requires a relatively low (Frisch) labor supply elasticity of 0.11 and substantial frictions that limit labor supply adjustment. On the normative side, these responses imply significant welfare gains from income-contingent repayment. Specifically, a constrained-optimal income-contingent loan provides gains relative to a 25-year fixed repayment contract equivalent to 1.3% of lifetime consumption, at the same fiscal cost. Equity contracts generate gains that are larger on average but more dispersed, suggesting they are less robust to unmodeled ex-ante responses. Finally, the moral hazard created

by income-contingent repayment decreases the insurance these contracts can provide at a given cost. In sum, my results suggest that income-contingent repayment creates (ex-post) moral hazard that affects contract design, but this moral hazard is too small to justify fixed repayment contracts.

There are several benefits to studying how income-contingent repayment affects labor supply in Australia. First, Australia was the first country to introduce income-contingent loans in 1989, meaning borrowers are familiar with the availability and design of these contracts, unlike in the US (Abraham, Filiz-Ozbay, Ozbay, and Turner 2020; Mueller and Yannelis 2021). Second, there is limited scope for adverse selection due to a lack of alternative financing options. This is useful for identification because it implies responses to the policy change reflect moral hazard rather than selection (Karlan and Zinman 2009), in contrast with the US, where lower-income borrowers select into income-contingent repayment (Karamcheva, Perry, and Yannelis 2020). Third, these loans can only cover tuition, which is mostly government-controlled, implying individuals can only adjust their borrowing by changing their degree choices. This decision is likely less responsive than the other margins borrowers in the US can adjust, such as room and board or groceries.

Figure 1. Income Distribution of Debtholders around Income-Contingent Loan Repayment Threshold



Notes: This figure shows the distribution in Australian dollars of the income that determines individuals' repayment rates on their income-contingent loan in 2004 and 2005 before and after the policy change. This income is called HELP Income and is equal to taxable income (i.e., the sum of labor income, capital income, and deductions), plus investment losses, retirement contributions, foreign employment income, and fringe benefits. The vertical lines indicate the threshold above which individuals begin making payments of 3% and 4% of their income in 2004 and 2005. The sample is the population of debtholders in Australia, subject to the sample selection criteria in Section 2.4. HELP Income is deflated to 2005 using the HELP indexation rate calculated from the Consumer Price Index.

I begin by documenting evidence of moral hazard from income-contingent repayment: individuals reduce their labor supply to minimize repayments on income-contingent loans. Figure 1 summarizes this behavioral response by plotting the income distribution of student debtholders in the two years surrounding the policy change. The vertical lines indicate the income-contingent threshold at which loan repayment begins, which was increased after the policy change. The income distribution exhibits significant bunching below the threshold both before and after the policy change, which I show is also present in the distribution of labor income. I present two pieces of evidence that suggest the bunching in Figure 1 reflects labor supply responses rather than solely

income-shifting or tax evasion. First, the bunching is larger in occupations with high hourly flexibility (e.g., bartenders) and almost non-existent in those with low flexibility (e.g., software engineers). Second, using data from Australia's Census, I find that individuals below the repayment threshold work 2-3% fewer hours (i.e., 1-2 fewer weeks per year) than those above the threshold.

Next, I develop a structural model of labor supply that quantitatively replicates the evidence in [Figure 1](#). The purpose of developing this model is to translate this evidence into estimates of preference parameters and study the normative implications of income-contingent repayment. In the model, overlapping generations of individuals choose consumption and labor supply over their life cycles. During working life, individuals repay their government-provided loans, and their labor income equals the product of endogenous labor supply and exogenous wage rates, where the latter is subject to uninsurable idiosyncratic risk. The two key ingredients in this model are uninsurable income risk and endogenous labor supply, which create a trade-off between the insurance benefits and moral hazard costs of income-contingent repayment.

The evidence in [Figure 1](#) is inconsistent with a frictionless formulation of this model in which labor supply is chosen to equate the marginal cost of working with the marginal benefits of higher income. When individuals' income crosses the repayment threshold, the fraction of *total* annual income they repay increases from 0% to 3-4%, or \$1,400 AUD in 2005 (\$1,800 USD in 2023). Under the standard assumption that utility is increasing in consumption and leisure, this model predicts no individuals would locate immediately above the threshold because locating below it delivers more leisure and \$1,400 more cash on hand.

Motivated by the evidence that labor supply responses increase with hourly flexibility, I add optimization frictions ([Chetty 2012](#)) to the model to explain individuals located above the repayment threshold. Because isolating the importance of every possible optimization friction is not feasible, I introduce two frictions that jointly characterize how several could affect labor supply adjustment in reduced form. First, in each period, only a fraction of individuals receive shocks that allow them to adjust labor supply à la [Calvo \(1983\)](#). These shocks could capture inattention or the arrival of job transitions at which hours can be adjusted. Second, adjusting labor supply requires paying a fixed cost, which could be monetary (e.g., wage reduction) or psychological (e.g., hassle costs).

I estimate the model by conducting the policy change from [Figure 1](#) in the model and find that, although the responses in [Figure 1](#) may appear large, rationalizing them requires labor supply to be relatively inelastic. The model's key parameters for determining labor supply responses are the (Frisch) labor supply elasticity, fixed adjustment cost, and Calvo probability. The labor supply elasticity is identified by the extent of bunching below the repayment threshold: a larger elasticity implies more bunching. The number of individuals above the repayment threshold then identifies the adjustment cost and Calvo probability: without these frictions, no individuals would be above the threshold. The estimation results show that replicating the evidence in [Figure 1](#) requires a labor supply elasticity of 0.11, fixed adjustment cost of \$400 (i.e., 1% of mean earnings), and Calvo

probability of 0.2. This estimate of the labor supply elasticity is similar to the mean elasticity of 0.15 from the meta-analysis in [Chetty \(2012\)](#). However, this meta-analysis is among studies without optimization frictions, while the frictions I estimate are quantitatively large: in a misspecified model without both frictions, the estimated labor supply elasticity is 0.005.

The estimated model highlights two important drivers of individuals' labor supply responses that receive empirical support: borrowing constraints and the expectation of future debt repayment. Borrowing constraints increase the incentives to reduce labor supply by raising the value of the additional cash on hand from locating below the repayment threshold. In a counterfactual where individuals can freely borrow at the riskless rate, the model predicts the bunching in [Figure 1](#) disappears almost entirely. Empirically, this importance of liquidity is supported by the fact that individuals below the repayment threshold have larger housing payments, which represent greater liquidity demands, are younger, which is when borrowing constraints are tightest, and have less wealth in the form of retirement savings. The second important driver of labor supply responses is that debt repayment ceases after initial balances are repaid: the bunching in [Figure 1](#) is almost twice as large in a counterfactual where repayments continue indefinitely. This finding is consistent with the fact that bunching increases with debt balances and highlights an important way income-contingent loans' effects on labor supply differ from those of income taxes.

In the final part of the paper, I use my structural model to study contract design and find that contracts with income-contingent repayment provide welfare gains relative to standard debt (i.e., fixed repayment) contracts, even after considering the moral hazard they create. My analysis considers a social planner that maximizes borrowers' lifetime utility by choosing one mandatory repayment contract, holding fixed borrowing behavior. This perspective isolates the central trade-off in income-contingent repayment between providing insurance and disincentivizing labor supply.

My main normative result is that income-contingent loans can simultaneously generate meaningful welfare gains and identical fiscal costs to fixed repayment contracts. I consider income-contingent loans with two parameters, as in the US: an income threshold at which repayment begins and a repayment rate of income above this threshold. I then solve for the values of these parameters that maximize borrower welfare subject to the constraint of raising the same revenue as a fixed repayment contract. The resulting constrained-optimal income-contingent loan provides gains equivalent to 1.3% of lifetime consumption relative to a 25-year fixed repayment contract, which is currently offered in the US and has a similar repayment duration without income-contingent payments. The cost of the moral hazard from income-contingent repayment is small: the consumption-equivalent gain from an alternative (infeasible) contract with wage-contingent repayments, which provides insurance without distorting labor supply, is only 0.2pp higher at 1.5%. Despite this small cost, labor supply responses quantitatively affect contract design: if labor supply did not respond to income-contingent repayment, the optimal contract would provide more insurance to low-income borrowers with a 40% higher repayment threshold.

I conclude by studying the welfare impact of three alternative income-contingent repayment contracts: income-contingent loans with forgiveness, fixed repayment contracts with forbearance, and income-sharing agreements. First, adding forgiveness to income-contingent loans after a fixed horizon, as done in the US and UK, generates losses relative to the constrained-optimal income-contingent loan. For a given fiscal cost, forgiveness increases repayments for young relative to old borrowers, which leads to losses because young individuals have a higher marginal value of wealth. Second, a fixed repayment contract with forbearance, a form of income-contingency that pauses payments for low-income borrowers, generates losses relative to the constrained-optimal income-contingent loan. This is because income-contingent loans accelerate payments from high-income borrowers, enabling them to provide more insurance at a given cost. Finally, equity contracts known as income-sharing agreements, which were originated by [Friedman \(1955\)](#) and recently implemented by Purdue University, yield gains that are larger on average but significantly more dispersed than those of income-contingent loans. This finding suggests that equity contracts cause ex-ante responses not captured by the model, implying that income-contingent loans may be a more robust mechanism for implementing income-contingent repayment.

Related literature and contribution. This paper sits at the intersection of literatures in household finance, public finance, and macro-finance. In its focus on the trade-off between insurance and moral hazard, this paper is part of a large literature on various forms of social insurance ([Chetty and Finkelstein 2013](#)), such as unemployment insurance ([Gruber 1997](#); [Ganong and Noel 2019](#)), debt relief ([Ganong and Noel 2020](#)), bankruptcy protection ([Dobbie and Song 2015](#); [Auclert, Dobbie, and Goldsmith-Pinkham 2019](#); [Indarte 2023](#)), and health insurance ([Einav, Finkelstein, and Schrimpf 2017](#)). Two strands of this literature that focus on student debt are directly related ([Amromin and Eberly 2016](#)). The first documents forms of debt overhang, in which reductions in student debt decrease delinquencies ([Di Maggio, Kalda, and Yao 2021](#)), increase homeownership ([Mezza, Ringo, Sherlund, and Sommer 2020](#)), and change choices of majors, degrees, and occupations ([Luo and Mongey 2019](#); [Chakrabarti, Fos, Liberman, and Yannelis 2020](#); [Folch and Mazzone 2021](#); [Ji 2021](#); [Hampole 2022](#); [Murto 2022](#); [Huang 2022](#)).¹ The second strand studies income-contingent loans as a tool to mitigate these effects, finding reductions in unsecured delinquencies ([Herbst 2023](#)), mortgage defaults ([Mueller and Yannelis 2019](#)), and the passthrough of income variation to consumption ([Gervais, Liu, and Lochner 2022](#)).² Taken together, this evidence provides empirical support for using equity-like contracts to finance human capital.

This paper makes three contributions to these literatures. First, it empirically identifies ex-post moral hazard (i.e., labor supply responses) created by income-contingent repayment, which has not

¹A related literature in labor economics emphasizes the importance in credit constraints for college attendance ([Carneiro and Heckman 2002](#); [Belley and Lochner 2007](#)). Compared to a counterfactual with no loans, student loans can help relax credit constraints and increase degree completion ([Black, Denning, Dettling, Goodman, and Turner 2022](#)).

²Alternative options to providing insurance are to make student debt dischargeable, which has the cost of inducing strategic default ([Yannelis 2020](#)), universal loan forgiveness, which would be regressive ([Catherine and Yannelis 2023](#)), and targeted loan forgiveness, which borrowers appear to value but fail to take-up ([Jacob, Jones, and Keys 2023](#)).

been found in other settings (Britton and Gruber 2020).³ Second, this paper provides a dynamic model of labor supply that rationalizes these responses, finding an important role for liquidity constraints and optimization frictions. Finally, it quantifies what these responses imply for the design of constrained-optimal repayment contracts. Prior literature highlights the insurance benefits of income-contingent loans, but has not had evidence to discipline their moral hazard effects or characterized optimal policy (Ji 2021; Matsuda and Mazur 2022; Boutros, Clara, and Gomes 2022).

This paper is also related to the literature on human capital financing. The idea that student loans should be equity-like was popularized by Friedman (1955), who advocated the use of income-sharing agreements. Adverse selection prevents the private provision of such contracts (Herbst and Hendren 2021; Herbst, Palacios, and Yannelis 2023), so a growing number of governments have attempted to correct this market failure by introducing income-contingent loans, with Australia being the leading example (Chapman 2006) and other countries following (Barr, Chapman, Dear-den, and Dynarski 2019). Theoretical work suggests these loans provide a close approximation to optimal policies (Lochner and Monge-Naranjo 2016; Stantcheva 2017).⁴ This paper contributes by quantifying how the moral hazard these loans create affects optimal contract design.

By studying state-contingent contracts, this paper is part of the literature on household security design. Motivated by evidence of imperfect risk-sharing (Cochrane 1991) and the household balance sheet channel (Mian and Sufi 2014), this literature studies policies that make liabilities more state-contingent, such as shared-appreciation mortgages (Caplin, Carr, Pollock, and Tong 2007; Hartman-Glaser and Hébert 2020; Greenwald, Landvoigt, and Van Nieuwerburgh 2021) or adjustment-rate mortgages conditioned on aggregate shocks (Campbell, Clara, and Cocco 2021). This paper contributes by studying one of the longest-running examples of such policies and characterizing the welfare gains from alternative forms of state-contingent repayment. An important feature of my setting is limited strategic default, as student loans cannot be discharged in bankruptcy.

Finally, this paper builds on extensive literature that exploits non-linearities in tax schedules to identify labor supply elasticities (Saez 2010; Chetty, Friedman, Olsen, and Pistaferri 2011; Kleven and Waseem 2013) by making two contributions. First, it empirically characterizes how labor supply responds to income-contingent repayment, which creates dynamic incentives that taxes do not. Second, it uses these responses to estimate a dynamic model with fixed costs and Calvo adjustment, which have been used to model optimization frictions but not separately estimated (Werquin 2015). The labor supply elasticity I estimate is on the low-end of estimates in this literature. This could reflect the fact that my sample is close-to-median income debtholders, who are early in their life cycles and may have less hourly flexibility or willingness to make extensive margin adjustments.

³This paper builds on Chapman and Leigh (2009), who study the Australian student loan system using survey data.

⁴Other parts of this literature study other government policies towards human capital, such as subsidies for educational expenses (Benabou 2002; Bovenberg and Jacobs 2005), expanding access to grants (Abbott, Gallipoli, Meghir, and Violante 2019; Ebrahimian 2020), and relaxing parental credit constraints (Caucutt and Lochner 2020).

1 Motivating Framework

This section develops a simple framework to clarify the trade-off between insurance and incentives created by income-contingent repayment. The result is an expression that generalizes the Baily-Chetty formula (Baily 1978; Chetty 2006) for the optimal balance of insurance and incentives in unemployment insurance to my setting. I then discuss the behavioral responses I attempt to estimate empirically through the lens of this expression.

Environment. Consider a government who provides a student loan, D_0 , at $t = 0$ to an individual in exchange for mandatory repayments $d_t = d(D_t, y_t, \theta)$ for $t > 0$, where D_t denotes the outstanding debt balance, y_t denotes observable income, and θ are the parameters of a repayment contract. For example, an equity contract is captured by $d_t = y_t \theta$, while a debt contract would be a function of just D_t and θ . Individuals solve a standard life cycle problem by choosing labor supply, ℓ_t , consumption, c_t , and initial debt balances, D_0 :

$$V(\theta) = \max_{\{c_t, \ell_t\}_{t=0}^T, D_0} \mathbf{E}_0 \sum_{t=0}^T u^t(c_t, \ell_t),$$

$$c_t + A_{t+1} = A_t R + y_t - d_t * \mathbf{1}_{t>0} + D_0 * \mathbf{1}_{t=0},$$

$$y_t = f(\ell_t, D_t, \omega_t), \quad d_t = d(y_t, \theta), \quad D_{t+1} = D_t R_d - d_t.$$

Expectations are taken over the path of stochastic shocks, $\{\omega_t\}_{t=0}^T$, which present income risk to the individual and are not observable to the government (Mirrlees 1974). Individuals can only take the government-provided contract and have no other sources of external financing.

Planner's problem. The government chooses θ to maximize borrower welfare. Assume all individuals are ex-ante identical so the government solves the following problem:

$$\max_{\theta} V(\theta) - \lambda' \left[D_0 - \sum_{t=0}^T \frac{E_0(d_t)}{\mathcal{R}_t} \right], \quad (1)$$

where λ' denotes the marginal cost of public funds, or equivalently the multiplier on the government budget constraint, and \mathcal{R}_t denotes the government discount rate at horizon t . The following proposition characterizes the optimal financing contract via a perturbation, as in Saez (2002).

Proposition 1. Define $M_t = \frac{u_c^t(c_t, \ell_t)}{u_c^{0t}(c_0, \ell_0)}$ as individuals' time t to time 0 stochastic discount factor, $\lambda = \lambda' \frac{\partial A_0}{\partial V}$ as the marginal cost of public funds in dollars, and θ^* as a solution to (1). Under appropriate

regularity conditions, the following condition holds at $\theta = \theta^*$:

$$\sum_{t=1}^T E_0 \left[\underbrace{\left(\frac{\lambda}{\mathcal{R}_t} - M_t \right) \frac{\partial d_t}{\partial \theta}}_{\text{amount of risk-sharing}} \right] = \lambda \left[\underbrace{\frac{dD_0}{d\theta}}_{\text{borrowing response}} - \sum_{t=1}^T \frac{1}{\mathcal{R}_t} E_0 \left(\underbrace{\frac{\partial d_t}{\partial y_t} \frac{dy_t}{d\theta}}_{\text{labor supply response}} \right) \right]. \quad (2)$$

The left-hand side of (2) is the *quantity of unshared risk*: it represents the difference between how the government values a perturbation to the repayment contract, $\frac{\partial d_t}{\partial \theta}$, and how the individual values it. If the government fully insures the individual, then the individuals' stochastic discount factor does not vary across states (for a given t), and this quantity is small. In contrast, if the individual is not fully-insured, then the difference between these valuations is large. The right-hand side of (2) the sum of two behavioral responses. The first is an *ex-ante* moral hazard effect, $\frac{dD_0}{d\theta}$: changing the repayment contract affects how much individuals borrow. The second behavioral response represents *ex-post* moral hazard: changing the repayment contract affects individuals' incentives to adjust their income, which affects the amount the government collects in repayments.

As an example, consider a policy change $d\theta$ that increases the amount of insurance by making low-income individuals pay less and high-income individuals pay more. In response, a natural prediction is risk-averse individuals will borrow more ex-ante, $\frac{dD_0}{d\theta} > 0$, and low-income individuals will increase their labor supply, $\frac{dy_t}{d\theta} > 0$, and high-income individuals will reduce their labor supply, $\frac{dy_t}{d\theta} < 0$. The heart of the insurance-incentive trade-off is illustrated in (2): if these responses are small, the government can afford to bear most of the income risk. If they are large, individuals must bear most of the risk to limit borrowing and encourage labor supply.

The objective of this paper is to quantify the magnitude of *ex-post* moral hazard in income-contingent repayment, $\frac{dy_t}{d\theta}$, and study what it implies for optimal contract design. To do so, I leverage a setting with a change in the repayment contract, $d\theta$, that allows me to estimate $\frac{dy_t}{d\theta}$. One of the main benefits of this setting (discussed in Section 2) is that individuals have limited ability to adjust initial debt balances, which reduces the scope for ex-ante moral hazard, $\frac{dD_0}{d\theta}$.

2 Institutional Background and Data

2.1 Overview of Australia's Higher Education Loan Programme (HELP)

In Australia, higher education is primarily financed using government-provided student loans through the Higher Education Loan Programme (HELP). HELP was introduced in 1989 and before 2005 was called the Higher Education Contribution Scheme (HECS). HELP loans can be used to

finance tuition⁵ for all undergraduate and graduate degree programs at public institutions.⁶ Tuition at public institutions is controlled by the government and varies by discipline. For undergraduate courses at these institutions, which are called Commonwealth Supported Places (CSPs), the government provides a subsidy in the form of a contribution to the tuition owed by the student. The remaining tuition after deducting the government’s contribution is paid by the student and called the student contribution. As of 2023, student contributions ranged from \$4,124 to \$15,142 AUD per year (\$2,700 to \$10,100 USD), and undergraduate degrees typically last 3-4 years. The number of CSPs in Australia has generally been capped by the government, except 2012-2017 in which the system was “demand-driven” (D’Souza 2018; Norton 2019).

Australian citizens can either pay their student contribution upfront or borrow through HELP. Most individuals choose to do the latter, with less than 10% of new debt in 2022 being paid upfront (Department of Education and Training 2023). If individuals borrow through HELP, their initial debt is equal to their student contribution.⁷ Given an average undergraduate student contribution of around \$6,000 USD per year, these debt burdens are comparable to tuition for US in-state public undergraduate programs, which averages around \$9,000 (Hanson 2023). Figure A1 plots the time series of student contributions for undergraduate degrees at public institutions, which varies based on whether the degree belongs to one of three bands, and the aggregate amount of HELP borrowing and upfront payments.

HELP debt balances in subsequent years grow at the CPI inflation rate net of repayments, meaning HELP debt has a zero real interest rate. Individual i ’s annual compulsory repayment in a given year is equal to

$$\text{HELP Repayment}_{it} = \min\{r_t(y_{it}) * y_{it}, D_{it}\},$$

where y_{it} denotes HELP Income, $r_t(\cdot)$ is the income-dependent repayment rate, and D_{it} denotes the current debt balance. HELP Income, y_{it} , is the taxable income reported in a personal income tax return plus a few adjustments discussed in Section 2.5. The collection of HELP payments is integrated with the personal income tax system and all individuals file tax returns in Australia, so y_{it} refers to *individual* rather than household HELP Income. For most individuals, HELP repayments are withheld by their employer throughout the year and deducted from their debt balances after filing their tax returns. Individuals also have the option to make voluntary repayments at any time. Additional details on the timing and structure of repayments is presented in Appendix C.

⁵To finance non-tuition expenses, students on income support can use a **Student Start-Up Loan**, but these loans only supported less than 100,000 borrowers in 2020-21. All other students must self-finance these expenses, which is generally done using credit cards or taking jobs.

⁶Private institutions play a relatively small role in Australia, making up only 3 out of 43 universities and 5.8% of the domestic enrollment share as of 2021. These institutions are slightly more popular among international students, with 11.7% of the enrollment share. Private institutions are much more expensive than public ones, especially for domestic students, and primarily compete by offering more niche products.

⁷Universities receive the student contribution regardless of whether individuals borrow through HELP because, in the latter case, it is paid by the Australian Government.

Repayment of HELP debt continues until the remaining balance equals zero or the time of death. Partial repayment is common: as of 2004, around 25% of debt balances were forecasted to be written off due to death, and in 2019 that estimate was 36% (Martin 2004; Robinson 2019). This means HELP effectively forgives debt for borrowers at the end of their working life when they stop generating sufficient income to make compulsory repayments, similar to the forgiveness embedded in US income-driven repayment plans. As in the US, HELP debt cannot be discharged in bankruptcy.⁸

2.2 2004-2005 Policy Change to HELP Repayment Rates

The policy change I exploit is a 2004-2005 change in the HELP repayment rate function, $r_t(\cdot)$. Figure 2 plots repayment rates as a function of real HELP Income prior to the policy change in blue and after the policy change in red. The most significant change was the location of the repayment threshold, which is the point at which individuals have to start making repayments, from around \$26,000 AUD to \$35,000 AUD. The average HELP Income among debtholders is around \$30,000, so this policy change generated reductions in required repayments for many individuals. It also generated an increase in repayment rates for high-earners with incomes above \$50,000. This policy change applied to all new and existing HELP debtholders.

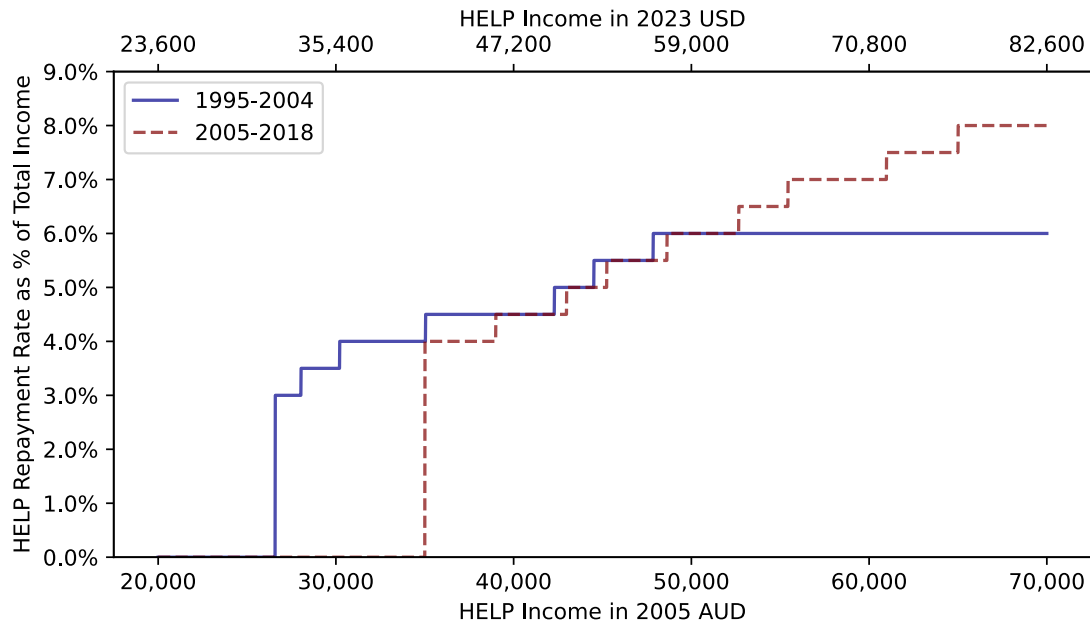
The repayment threshold creates a large incentive to reduce HELP Income by generating a discontinuity in the *average* rather than marginal repayment rate. For example, consider an individual with \$35,000 of HELP Income in 2005. For this individual, earning an extra \$1 of income results in a required HELP repayment of $\$35,001 \times 4\% \approx \$1,400$. In a frictionless static model of labor supply, no individuals would choose to locate immediately above the repayment threshold because doing so delivers less take-home pay and leisure relative to locating below it (i.e., it is a notch in the language of Kleven and Waseem 2013). The repayment functions feature several other discontinuities in average repayment rates shown in Figure 2. These are smaller: 0.5% versus 3%-4% at the repayment threshold in 2004 and 2005.

XThe fact that I find liquidity constraints, which increase the value of this repayment reduction, play a significant role in labor supply responses is consistent with the importance of liquidity in driving responses to other social insurance programs (Chetty 2008; Ganong and Noel 2022; Indarte 2023).

There are several reasons to believe this change to the HELP repayment function was salient to debtholders. First, the policy received media coverage at the time of the change (Marshall 2003). Second, the repayment function is indexed to inflation, which means it updates every year. When it is published at the beginning of each tax year, the government makes that receives press coverage

⁸Aside from death, the only case in which HELP debt is canceled is if an individual withdraws from the corresponding units of study before the census date in a given year.

Figure 2. HELP Repayment Rates as a Function of Income: Before and After 2004-05 Policy Change



Notes: This figure shows HELP repayment rates as a percentage of HELP Income for different levels of HELP Income. The blue line shows repayment rates in 2004 before the policy change, and the red line shows repayment rates in 2005 after the policy change. See Figure A2 for a plot of how marginal HELP repayment rates vary with income.

(Medhora 2018).⁹ Finally, the fact that HELP Income determines repayment rates with a repayment threshold has not changed since the introduction of HECS in 1989, meaning debtholders are likely to understand the program's structure.

Government policy documents and media articles suggest the primary reason for this change was to reduce the burden placed on lower-income individuals, for whom payments were burdensome and contributed little to the total HELP budget (Nelson 2003; Marshall 2003). In addition to changing the HELP repayment function, other policy changes were implemented in 2004-2005, such as the introduction of HELP loans for private undergraduate and graduate degrees through FEE-HELP and increase of student contributions by 25% (see Figure A1). These other changes, reviewed by Chapman and Salvage (2001) and Beer and Chapman (2004), were primarily aimed at those entering their degree programs rather than those repaying HELP debt. The simultaneous implementation of these other changes with the change in repayment threshold is not ideal. However, it likely has a minimal effect on my analysis, which focuses on identifying moral hazard among individuals who have already completed their degree programs.

⁹For an example of an announcement, see <https://www.legislation.gov.au/Details/C2022G00213>.

2.3 Benefits of Studying Income-Contingent Repayment in Australia

In addition to high-quality administrative data and policy variation, there are several benefits to using HELP to identify labor supply responses to income-contingent repayment. First, there is limited room for adverse selection of high-income individuals into non-income-based repayment contracts because HELP is the only government-provided student loan, which implies the effects of changes in contract design reflect borrowers' actions rather than types (Karlan and Zinman 2009).¹⁰ The same is not true in the US, where high-income borrowers choose fixed rather than income-driven repayment (Karamcheva et al. 2020), nor in countries with private providers of income-sharing agreements (Herbst et al. 2023). In principle, individuals in Australia could receive external financing from a bank or university. However, there is little economic incentive to do this because the interest rate would exceed the zero real interest rate on HELP. The primary margin in which there is scope for adverse selection is on whether to pay upfront or borrow through HELP, but the zero interest rate on HELP again implies little incentive to pay upfront. In practice, the amount of upfront borrowing has been low and stable, with most payments coming from individuals with family support (Norton 2018).

A second benefit of this setting is likely limited *ex-ante* moral hazard, in which individuals increase (decrease) their initial HELP debt in anticipation of a lower (higher) probability of future repayment. As described above, HELP can only be used to cover tuition and tuition at public undergraduate institutions, which make up over 94% of the domestic enrollment share and 39/42 universities is controlled by the government. As a result, individuals can only adjust their initial HELP debt by changing their choice of degree or institution, which are likely stickier decisions than the other margins borrowers in the US can adjust, such as room and board or groceries.

The third benefit of studying HELP is that it is the longest-running government-provided income-contingent repayment program. The fact that this program has been around since 1989 suggests that borrowers understand the structure of income-contingent repayments. The same is not true in the US, where borrowers are unaware of the existence and structure of income-driven repayment options (Abraham et al. 2020; Mueller and Yannelis 2021; JPMorgan Chase 2022). A final benefit is that there are no responses by the supply-side of higher education due to government tuition control. If this were not the case, changes in government-provided financing contracts could pass through to tuition and thus initial debt balances (Kargar and Mann 2022).

One caveat of using HELP to identify labor supply responses to income-contingent repayment is that this program is heavily-subsidized. A less subsidized program would only draw in individuals that place higher values on education. If there is heterogeneity in the structural parameters govern-

¹⁰This approach of identifying moral hazard by looking at the responses to changes in contract structure among individuals who have already taken up the contract has been applied in a variety of selection markets, such as consumer credit (Einav, Jenkins, and Levin 2012; DeFusco, Huan, and Yannelis 2022), mortgages (Gupta and Hansman 2022), and health insurance (Einav and Finkelstein 2011).

ing labor supply that is correlated with the value of education, a program with a different subsidy would generate different labor supply responses. To mitigate this concern, my counterfactual analysis focuses on repayment contracts with a similar fiscal cost to HELP.

2.4 Data Sources

I use several sources of restricted-access de-identified administrative data. First, I use individual income tax returns from the Australian Taxation Office (ATO), which contain panel data on income components and basic demographic characteristics. Second, I use administrative data on HELP from the ATO that includes debt balances, repayments, and a flag for whether individuals acquired new debt balances in a given year. Third, I leverage administrative data on superannuation balances and contributions from the ATO. These three datasets are linked for the universe of Australian taxpayers between 1991 and 2019 in the [ATO Longitudinal Information Files](#), known as *ALife*. Starting from the population dataset in *ALife*, I restrict attention to individual-year observations that are (i) between ages 20 and 64, (ii) residents in Australia for tax purposes, (iii) not exempt from HELP repayment due to a Medicare exemption, and (iv) do not have any income from discretionary trusts.¹¹ I use this sample for my main analysis that only requires tax and HELP data.

To obtain data on hours worked and housing payments, I use a linkage of these ATO data with the 2016 Census of Population and Housing. This linkage cannot be performed with the *ALife* data directly, so I instead perform the merge through the [Australian Bureau of Statistics Multi-Agency Data Integration Project](#) (MADIP). The ATO data in MADIP has the same sample coverage as the population-level *ALife* data but a restricted set of variables. Due to data limitations, I use the first three filters from the *ALife* sample to construct a cross-sectional MADIP sample in 2016, the year in which the Census was administered.

I supplement these administrative datasets with survey data on household balance sheets from the [Household, Income and Labour Dynamics in Australia Survey](#) (HILDA), which is a household survey conducted by the Melbourne Institute that runs from 2002 to 2021. HILDA has a similar structure and questions to the [Survey of Consumer Finances](#) in the US, except that it is a panel rather than a repeated cross-section.

¹¹In Australia, there are unit trusts, in which trust beneficiaries have no discretion over entitlements, and discretionary trusts, in which beneficiaries have full discretion over entitlements. These two make up around 48% and 52% of trusts, respectively. There are two main sources of trust income: (i) investment income from listed and unlisted securities, including money market funds, equities, and rental properties (e.g., Vanguard-managed trust funds), and (ii) capital or labor income from businesses held in trust structure. These make up around 65% and 35% of total trust income, respectively. Discretionary trusts have been identified as potential sources of tax evasion ([Australian Council of Social Service 2017](#)), but *ALife* does not have information on the sources of trust income. Therefore, I choose to drop these observations from my analysis to avoid attributing possible tax evasion to labor supply responses. This represents less than 2.5% of debtholders at the time of the policy change in 2004. In results available upon request, I show my main findings are quantitatively similar when keeping these observations.

2.5 Summary Statistics

Table 1 presents summary statistics on the *ALife* sample, which is the main sample in my analysis. The three columns separate the sample into individuals without HELP debt, individuals with HELP debt, and 26-year-old HELP debtholders, which is the age at which most individuals have finished university in Australia and begun work, and average HELP debt balances peak in real terms. Relative to non-debtholders, debtholders tend to be younger, less likely to be wage-earners, which is defined as having any self-employment income from partnerships, sole-traders, or personal-services, and have lower taxable income.

Table 1. Summary Statistics

	Sample of Individuals		
	Non-Debtholders (1)	Debtholders (2)	26-Year-Old Debtholders (3)
Demographics			
Age	41.1	29.5	26
Female	0.46	0.60	0.57
Wage-Earner	0.85	0.91	0.93
Income Totals			
Taxable Income	37,695	27,796	32,929
HELP Income	38,756	28,586	33,721
Income Components			
Salary & Wages	32,415	26,068	32,091
Labor Income	35,480	27,136	32,999
Interest & Dividend Income	726	242	224
Capital Income	1,221	324	184
Net Deductions	-1,548	-1,099	-554
HELP Variables			
HELP Debt	.	10,830	13,156
HELP Payment	.	991	1,305
HELP Income < 0% Threshold	0.50	0.65	0.51
HELP Income < 2004 0% Threshold	0.37	0.51	0.35
HELP Income < 2005 0% Threshold	0.52	0.67	0.55
Number of Observations	247,118,713	27,316,037	1,701,464

Notes: This table presents summary statistics from the *ALife* sample from 1991-2019, subject to the sample selection criteria discussed in Section 2.4. The values for all continuous variables represent means. All continuous variables are deflated to 2005 dollars based on the HELP threshold indexation rate. All continuous variables except HELP Debt and HELP Repayment are winsorized at 2%-98%. HELP Income < 0% Threshold corresponds to the mean of a dummy variable for whether HELP Income in an individual-year was below the 0% HELP repayment threshold. HELP Income < 0% 2004 Threshold and HELP Income < 0% 2005 Threshold correspond to means between 1998-2004 and 2005-2018 for whether HELP Income in an individual-year was below the HELP repayment threshold, respectively, after adjusting the thresholds for inflation. Additional details on variable construction are presented in Appendix D.

The most important variable introduced in **Table 1** is HELP Income, which determines an individual's repayment rate on their HELP debt according to **Figure 2**. HELP Income equals taxable income plus several other adjustments, such as adding back reportable superannuation contributions and investment losses. These adjustments are not relevant for most individuals: the difference between

HELP and taxable income in 2004 is less than \$100 for over 93% of observations in 2004.

I decompose HELP Income into three terms:

$$\text{HELP Income} = \text{Labor Income} + \text{Capital Income} - \text{Net Deductions.} \quad (3)$$

Labor Income is defined as the sum of salary and wages, tips and allowances, and self-employment income. This represents the largest source of income for most individuals: 95% for debtholders and 91% for non-debtholders. Capital Income is defined as the sum of interest income, dividend income, capital gains, government superannuation and annuity income, rental income, and trust income. Importantly, Capital Income does not capture flow income from owner-occupied housing, which cannot be inferred from income tax returns because Australia does not have a mortgage interest deduction. Net Deductions is defined as the residual in (3). Additional details on the construction of all variables are presented in Appendix D.

Table 1 shows debtholders have lower HELP Income, Labor Income, and Capital Income, in addition to fewer deductions, than non-debtholders. These differences are not surprising given the age differences between the two groups. The average debt balance among debtholders is around \$10,800 in 2005 AUD (\$13,000 in 2020 USD), and around \$13,200 in 2005 AUD (\$15,800 in 2020 USD) among 26-year-old debtholders. Notably, most debtholders (65%) in each year are below the HELP repayment threshold, especially after the 2004-2005 policy change. Focusing on 26-year-old debtholders who have likely finished college and entered the workforce, around half are below the threshold. The 2004-2005 policy change had a big impact for these individuals: the fraction below the threshold moved from 35% to 55% after the policy change.

Figure A3 shows how debt balances vary within-individual over time: most individuals' debt balances peak in real terms between ages 24 and 26, and are paid down in their mid-30s. However, around 15% of individuals who had debt at age 22 in 1991 still have debt at age 50 in 2019. Given the increase in real tuition over time, this number is forecasted to be much higher with around 36% of outstanding debt expected to be not paid off (Robinson 2019).

3 Labor Supply Responses to Income-Contingent Repayment

In this section, I use the discontinuities in repayment rates and the policy change to these rates, illustrated in Figure 2 to document several facts about how labor supply responds to income-contingent repayment.

3.1 Fact #1: Bunching Below Repayment Threshold

Figure 3 plots the distribution of real HELP income for individuals with HELP debt in the three years before and after the policy change. HELP Income is deflated to 2005 Australian dollars using the HELP Threshold indexation rate. The vertical line in each plot corresponds to the HELP repayment threshold in that year, which is constant in real terms across the years in which there is no policy change. In these plots, I focus on borrowers with HELP Income between \$20,000 and \$45,000 in 2005 dollars, which covers a substantial fraction of HELP debtholders. For example, the 2004 distribution in Figure 3 covers over 50% of the entire population of debtholders.

These results show significant bunching below the repayment threshold from 2002 to 2007. In the three years before the policy change, shown in the left three panels, the amount of bunching and shape of the income distribution remain relatively constant. However, after the policy change in 2005, the right three panels of Figure 3 show the income distribution exhibits two important changes. First, the bunching at the 2004 repayment threshold disappears completely. Second, bunching appears immediately below the new repayment threshold. This provides clear evidence that borrowers adjust their income to avoid making income-contingent repayments.

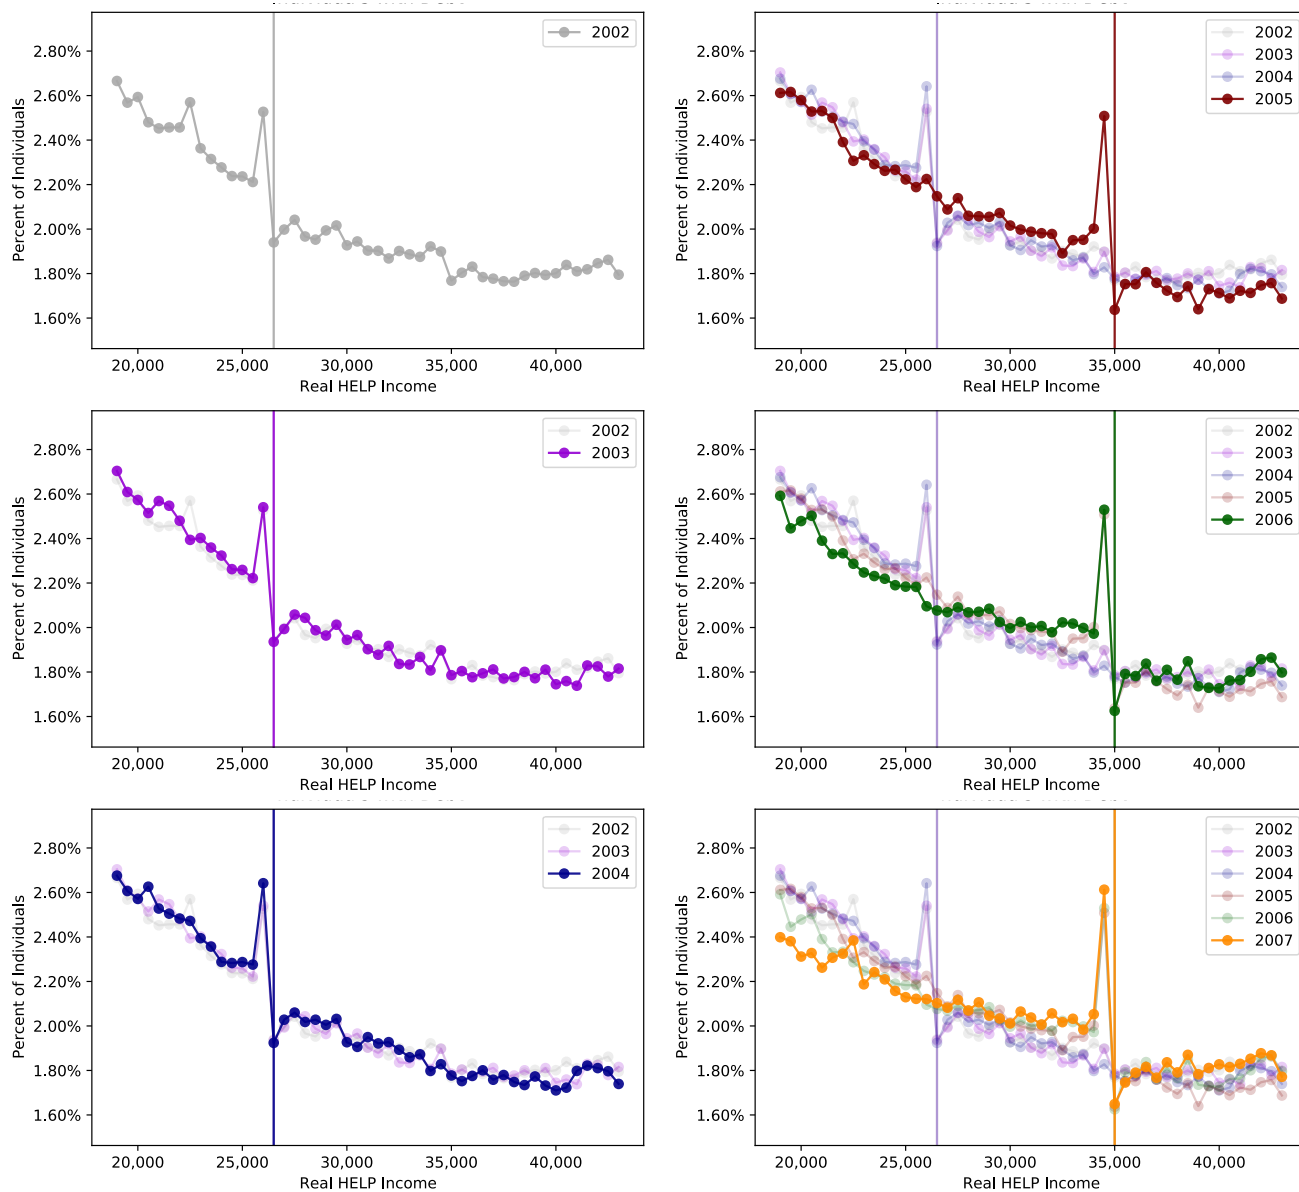
The fact that the bunching in Figure 3 responds quickly to the policy change shows it is not driven by mechanical features of Australia's tax system, such as the tendency to report incomes at round numbers. However, a possible threat to identification is the presence of other changes between 2002 and 2007 that affected individuals' incentives to report incomes of certain values. Although it is unlikely this could explain the evidence Figure 3, given the bunching is sharp around the repayment threshold, I assess this possibility by examining the income distribution of non-debtholders in Figure A4. In contrast to the income distribution of debtholders, these results show no change in the income distribution of non-debtholders around the repayment threshold both before and after the policy change.¹²

According to (3), HELP Income consists of three components that individuals could adjust to locate below the repayment threshold. Figure A5 shows that most of this response comes from Labor Income. In particular, I follow Chetty et al. (2011) and examine a sample of individuals whose primary source of income is Labor Income, earning less than 1% of income from other sources. This ensures that all individuals require similar values of Labor Income to generate HELP Income at the repayment threshold. I then compute a measure of bunching from Chetty et al. (2011) (described in Section 3.4) for the distributions of HELP and Labor Income. The results show the amount of bunching in Labor Income is 83% as large as that of HELP Income.¹³

¹²There are small changes in the income distribution of non-debtholders at lower values of income, which reflect changes in real terms of the second income tax bracket.

¹³The results in Figure A5 for HELP Income can be used to estimate the dollar loss to the ATO from the bunching at the repayment threshold: the HELP repayments implied by the counterfactual distribution for HELP Income estimated on the full sample from 2005-2018 are around \$90M higher than those implied by the observed distribution. This amounts to 42 bps of the total HELP compulsory repayments reported in the aggregated ATO HELP Data over this time period.

Figure 3. Real HELP Income Distribution of Debtholders Around Policy Change



3.2 Fact #2: More Bunching in Occupations with Greater Hourly Flexibility

Next, I explore variation in the bunching in [Figure 3](#) across different occupations. Using survey data on hours worked from HILDA, I compute a measure of the amount of hourly flexibility in each 2-digit ANZSCO occupation, which is the finest level at which *ALife* reports occupation codes, defined as the standard deviation of annual (within-individual) changes in (log) hours worked. This measure is highest for occupations where it is relatively easy to adjust hours, such as for Hospitality Workers (e.g., bartenders) and Food Preparation Assistants (e.g., fast-food workers), and lowest for those where it is more difficult, such as ICT Professionals (e.g., software engineers). [Table A1](#) shows these measures for each occupation code in my sample.

[Figure 4](#) plots the amount of bunching between 2005 and 2018 among wage-earners below the new repayment threshold relative to this measure of hourly flexibility. I focus on the period after the policy change because this is when *ALife* has comprehensive coverage of occupation codes. Each point represents a 2-digit occupation, and I measure the amount of bunching as the ratio of the number of individuals in that occupation within \$2,500 below to the number above the threshold, similar to [Chetty, Friedman, and Saez \(2013\)](#). A ratio of one indicates no bunching.

The results show that bunching is more common in occupations with greater hourly flexibility. For example, ICT Professionals have the lowest hourly flexibility with a standard deviation of annual change in log hours of 0.17. In this occupation, there is only 5% more individuals below relative to above the threshold. In contrast, Hospitality Workers have almost three times more hourly flexibility with a standard deviation of annual change in log hours of 0.48 and exhibit significantly more bunching, with 80% more individuals below relative to above the threshold. [Figure A6](#) shows a similar pattern using an alternative measure of hourly flexibility.

3.3 Fact #3: Borrowers Below Repayment Threshold Work Fewer Hours

To document the third fact, I use the responses from a question asked in the 2016 Census of Population and Housing in which individuals report the number of hours worked in all jobs during the week before Census night. [Figure 5](#) plots average hours worked in \$250 bins of HELP Income around the repayment threshold, in addition to the distribution of HELP Income in red. The results show that individuals locating immediately below the threshold work on average around 1 hour less per week than those immediately above the threshold. The standard work week in Australia is 38 hours, so this corresponds to a reduction of 2.6%.¹⁴ This adjustment in hours worked occurs *within* an individuals' current occupation: [Figure A8](#) finds little evidence that those below the repayment threshold are more likely to have switched occupations.

¹⁴These results are not driven by a group of individuals outside the labor force earning only income from other sources: [Figure A7](#) shows the patterns are nearly identical in the sample of individuals earning positive Labor Income.

Figure 4. Variation in Bunching across Occupations based on Hourly Flexibility





Figure 5. Average Hours Worked around Repayment Threshold



The results in [Figure 5](#) are subject to a few caveats. First, this test can only be performed in 2016 because this is the only year of Census data available in MADIP. Second, as discussed in [Section 2.4](#), the MADIP and *ALife* samples differ slightly. To mitigate concerns about sample selection, [Figure A9](#) shows the distribution of HELP Income in 2016 across the two samples is very similar. Finally, these data on hours worked are self-reported by employees rather than employers, which introduces concerns about reporting issues. As a result, I do not target this evidence directly when estimating my structural model but instead view it as providing qualitative support for the mechanisms included in the model.

3.4 Fact #4: Bunching Increases with Debt and Decreases with Age


The next two facts come from exploring heterogeneity in the bunching from [Figure 3](#) with observable characteristics. To measure the amount of bunching systematically, I construct a bunching statistic following the literature that uses discontinuities in tax rates to estimate taxable income elasticities ([Chetty et al. 2011](#); [Kleven and Waseem 2013](#)). First, I fit a five-piece spline to each distribution leaving out the region $\mathcal{R} = [\$32,500, \$35,000 + X]$. The choice of \$32,500 represents a conservative estimate of where the bunching begins and X is a constant intended to reach the upper bound at which the income distribution is affected by the threshold. This spline corresponds to an estimate of the counterfactual distribution absent the threshold. Next, I iterate on X so that this counterfactual density integrates to 1. Finally, I compute the bunching statistic, b , as:

$$b = \frac{\text{observed density in } \mathcal{R}}{\text{counterfactual density in } \mathcal{R}} - 1. \quad (4)$$

This bunching statistic is an estimate of the excess number of individuals below the repayment threshold relative to a counterfactual distribution in which the threshold did not exist. [Appendix E](#) presents additional details on this procedure.

[Figure 6](#) shows this estimated bunching statistic across groups of individuals with different ages and debt balances. I split debt balances at their median value within each year and then split ages into five year bins, which gives a similar number of observations within each bin. The results show two patterns. First, the amount of bunching increases in debt balances: at all age groups except 35-39, the estimated value of b is higher among individuals with above median debt balances. Second, the amount of bunching decreases with age. For example, the estimated bunching statistic is 22 – 33% lower among individuals above 40 than those below 25. Given that borrowing constraints are tightest among young individuals, this finding provides suggestive evidence that liquidity affects labor supply responses, which is examined in the next section.

Figure 6. Variation in Bunching by Debt Balances and Age



3.5 Fact #5: Bunching Decreases with Proxies for Liquidity


The fifth fact comes from examining how the bunching in [Figure 3](#) correlates with proxies for the extent to which individuals are liquidity-constrained. A demand for liquidity created by incomplete markets has been identified as an important driver of individuals' responses to various social insurance programs, such as unemployment insurance ([Chetty 2008](#)), mortgage default ([Ganong and Noel 2022](#)), and consumer bankruptcy ([Indarte 2023](#)). Absent direct measures of liquidity, this section uses several complementarity measures to assess its importance.

First, I use data on superannuation balances from *ALife*. Superannuation (“super”) represents the largest form of retirement savings in Australia and the second-largest source of household wealth ([Australian Council of Social Service 2018](#)). Contributions into a super account primarily come from mandatory employer and voluntary employee super contributions. Employee contributions, up to a limit, have generally been taxed at a rate lower than the personal income tax rate to incentivize saving. The left panel of [Figure 7](#) plots the estimated statistic, b , based on quartiles of superannuation balances defined within each year. The results show the amount of bunching is highest for individuals in the bottom quartile of superannuation balances, around twice as large as the top quartile. This evidence is consistent with these individuals being more liquidity-constrained and thus unwilling to use liquidity for tax-advantaged retirement saving.

My second piece of evidence leverages data on annual combined mortgage and rent payments from the 2016 Census using the MADIP sample. For most individuals, housing payments represent one of the largest sources of their liquidity demands. Therefore, if liquidity influences labor supply

responses, individuals below the repayment threshold should have larger housing payments or, equivalently, individuals with larger housing payments should be more likely to bunch below the repayment threshold. The right panel of [Figure 7](#) shows this pattern holds in the data: individuals immediately below the repayment threshold tend to have larger housing payments by around \$500-\$1000 (i.e., 3-6%).

Figure 7. Bunching, Retirement Wealth, and Housing Payments



A final, more speculative finding that points to the importance of liquidity is presented in [Figure A12](#), which plots the relationship between the amount of bunching and house prices from CoreLogic across geographic regions. Absent comprehensive data on wealth at the individual-level, house prices are a reasonable proxy for wealth because housing represents the largest form of household wealth in Australia ([Australian Council of Social Service 2018](#)). The results show that the amount of bunching is lower in regions with higher house prices, which tend to be metropolitan areas (e.g., Sydney), and that this relationship is unaffected by controlling for demographic and economic characteristics, such as population size and the unemployment rate.

3.6 Fact #6: Limited Evidence of Future Wage Reductions from Bunching

My final fact comes from exploring the dynamic effects of the bunching in [Figure 3](#). In models with learning-by-doing (also known as human capital accumulation or career effects) ([Keane and Wolpin 1997](#); [Imai and Keane 2004](#); [Best and Kleven 2012](#); [Keane and Rogerson 2015](#); [Makris and Pavan 2021](#)), the choice of current labor supply affects the stock of human capital and hence future wages. As a result, these models predict that the reduction in labor supply shown in [Figure 5](#) comes at the cost of lower future wages.¹⁵

¹⁵A related model of dynamic compensation is presented in [Kleven, Kreiner, Larsen, and Sogaard \(2023\)](#), where individuals' realized earnings only equals their true latent earnings (hours \times wages) at job transitions. I cannot test this hypothesis in my setting because I do not observe job transitions, but two facts suggest it is likely a small driver of the lack of responses. First, the same of individuals around the repayment threshold are relatively low-income, while [Kleven](#)

In the ideal experiment to identify the size of these future wage reductions, bunching would be randomly-assigned and I could compare the future wages of bunchers and non-bunchers. Absent this ideal experiment, [Figure 8](#) plots the average growth rate in Labor Income from year t to $t + h$ based on individuals' locations relative to the repayment threshold in year t . The results show individuals that bunch below the repayment threshold in year t experience lower income growth than those above the threshold in the subsequent year. However, this difference is small, only 1%, and disappears after three years. Although this evidence is clearly subject to concerns about selection into bunching, a natural form of selection would be that individuals with lower expected income growth would be more likely to bunch. In this case, the evidence in [Figure 8](#), which suggests relatively small wage reductions from bunching, would serve as an upper bound. Nevertheless, this evidence should not be interpreted as suggesting no learning-by-doing is present: larger labor supply reductions could create sizeable longer-horizon costs. Instead, it suggests the size of the responses created by income-contingent repayment are not large enough to create costs over the horizons I observe.

Figure 8. Future Labor Income Growth around Repayment Threshold



3.7 Possible Other Mechanisms for Bunching Below Repayment Threshold

This section discusses the possibility of other mechanisms through which individuals could reduce their income to locate below the repayment threshold.

Evasion. An obvious explanation for the bunching in [Figure 3](#) is evasion, in which individuals [et al. \(2023\)](#) find dynamic compensation plays an important role at the top of the income distribution. Second, [Figure A8](#) shows no discontinuity in the probability of occupation switching around the repayment threshold.

misreport their incomes. Although this is illegal and difficult to identify empirically, several facts, in addition to the direct evidence of a labor supply response in [Figure 5](#), suggest it cannot explain all of these responses. First, [Figure A15](#) replicates the analysis [Figure A5](#) replacing Labor Income with salary and wages and shows this distribution also exhibits substantial bunching around the repayment threshold. Bunching in the distribution of salary and wages is generally interpreted as evidence of hours worked responses (see e.g., [Chetty et al. 2013](#)) because the literature on tax evasion that uses random audits finds the majority of individual tax evasion comes from self-employment income, with an estimated non-compliance rate for third-party reported items, such as salary and wages, of less than 1% ([Slemrod 2019](#)). Second, [Figure A16](#) shows the amount of bunching only declines by 6% when restricting to the sample of Wage-Earners, who have substantially less flexibility in reporting their income. Third, [Figure A14](#) shows the amount of bunching is almost identical between individuals who file their tax returns electronically and non-electronically. When filing electronically, pure evasion is much more difficult because sources of Labor Income are often pre-filled by the employer, and, if they are not, the ATO directly compares what the individual reports with the employer's payment summary. Finally, the sample of individuals near the repayment threshold is around median income. This contrasts with the high-income sample of individuals for which prior literature has documented such shifting responses are prevalent ([Slemrod and Yitzhaki 2002](#)).

Nevertheless, it is likely that at least some of the responses in [Figure 3](#) reflect evasion rather than labor supply. If this is the case, the model I develop in [Section 4](#) will overestimate how much labor supply responds to income-contingent repayment. There are two ways this could affect my normative results. First, if individuals equate the marginal benefits of evasion (i.e., reduced repayments) with the marginal *social* cost of evasion, then whether the responses in HELP Income reflect labor supply or evasion is irrelevant, as long as my model can replicate them ([Feldstein 1999](#)). In the more likely case where the private and social costs of evasion are not equal, my results would overstate the welfare costs of moral hazard created by income-contingent repayment ([Chetty 2009](#)), reinforcing my qualitative conclusions.

Income-shifting across years. The repayment threshold incentivizes individuals to transfer income to the future if they anticipate being above the threshold later on. In practice, this could take the form of employees asking employers to delay some of their compensation. [Figure 8](#) shows this does not happen empirically: individuals below the repayment threshold in a given year do not have higher income in future years.

Firm responses. An alternative mechanism to the labor supply response in [Figure 5](#) would be a demand-side response, in which firms offer jobs with wages below the repayment threshold. [Chetty et al. \(2011\)](#) provides evidence of such a response by firms to reduce income tax rates in Denmark, in which the vast majority of private-sector jobs are covered by collective bargaining agreements. Two findings suggest this does not occur in my setting. First, the distribution of non-

debtholders, who compete in the same labor market, does not exhibit any bunching, as shown in [Figure 3](#). Second, [Figure A13](#) replicates Figure 9 from [Chetty et al. \(2011\)](#), which plots the distribution of Labor Income among individuals with positive Net Deductions. In [Chetty et al. \(2011\)](#), this distribution still exhibits bunching around the threshold at which marginal tax rates change because firms offer jobs with salaries below the threshold, even though this threshold does not apply to these individuals who claim deductions. In contrast, [Figure A13](#) shows this distribution exhibits no bunching in my setting.

3.8 Summary of Empirical Results and Implications for Structural Model

Summary of results. The previous sections present empirical facts about how labor supply responds to income-contingent repayment that can be summarized as follows. First, borrowers reduce their income in response to income-contingent repayment. These responses reflect, at least in part, labor supply responses rather than entirely tax evasion or income-sharing, as borrowers below the repayment threshold work fewer hours and tend to be in occupations with more flexibility. Second, the size of labor supply responses to income-contingent repayment vary cross-sectionally based on two forces. The first force is dynamics: borrowers with higher debt balances, for whom the repayment reduction is more likely a permanent reduction rather than simply a transfer over time, exhibit greater responses. The second force is liquidity: borrowers who are likely to be liquidity-constrained, for whom the value of the repayment reduction is most valuable, also exhibit larger responses. Finally, there is limited evidence of a dynamic cost associated with the reductions in labor supply that income-contingent repayment creates.

Implications for model. In [Section 4](#), I develop a structural model motivated by this empirical evidence. Consistent with the bunching below the repayment threshold and the importance of dynamics and liquidity, the model is a dynamic life cycle model in which individuals choose their labor supply by trading off the disutility of work with the benefits of higher income and choose consumption subject to borrowing constraints. However, the evidence in [Figure 3](#) also provides a rejection of a model in which labor supply is determined *solely* by trading off the disutility of work with the benefits of higher income. To see this, recall that locating immediately below the threshold increases take-home pay. If utility increases in consumption and leisure, such a model cannot generate any individuals immediately above the threshold because locating below it strictly dominates: it delivers more consumption with less labor supply ([Kleven and Waseem 2013](#)).¹⁶

The presence of individuals above the repayment threshold thus raises the question of what

¹⁶One reason individuals may locate above the repayment threshold is that, unlike a tax, income-contingent loans have an additional effect: increasing labor supply today reduces the stock of debt tomorrow. If the value function is sufficiently decreasing in debt, it may be optimal not to locate below the threshold. In [Appendix B](#), I show that for this to be the case, it must be the case that the value function is decreasing in debt more than it is increasing consumption. This is unlikely to be the case because HELP debt has a zero real interest rate, and individuals have the option to make voluntary repayments, but the majority do not.

mechanism should be added to the model to explain this lack of labor supply adjustment. Broadly speaking, there are three possible explanations. First, individuals may be unaware of the repayment threshold due to inattention (Chetty, Looney, and Kroft 2009; Chetty et al. 2013). Second, individuals may be aware of the threshold, but may be unable to adjust their labor supply due to costs associated with changing labor supply (Chetty 2012) or hours constraints (Chetty et al. 2011). Finally, individuals may be able to adjust their labor supply but actively choose not to locate below the repayment threshold. This could be because of long-run costs associated with doing so (Keane and Rogerson 2015), individuals receiving non-pecuniary benefits from work, or pro-social preferences in which individuals feel obligated to repay their debts. The model in Section 4 introduces optimization frictions to explain the presence of individuals above the repayment threshold, which capture the first two explanations but not the third. This choice is motivated by my finding that the amount of bunching increases with hourly flexibility, which suggests hours constraints and adjustment costs play a role, and the limited evidence of future wage reductions associated with bunching.

4 The Model

The empirical analysis in Section 3 documents labor supply responses to income-contingent repayment. However, this analysis leaves open two important questions. First, how large are these responses quantitatively? Second, are these responses large enough to imply the moral hazard costs created by income-contingent repayment outweigh its insurance benefits? The section presents and estimates a structural model designed to answer these two questions. The key ingredients in the model are endogenous labor supply, which creates moral hazard in response to income-contingent repayment, and uninsurable income risk, which creates a demand for insurance that income-contingent repayment can provide.

4.1 Model Description

4.1.1 Demographics

Time is discrete, and each period, t , corresponds to one calendar year. At time $t = h \in \{\underline{h}, \underline{h} + 1, \dots, \bar{h}\}$, a cohort h of individuals indexed by i are born at an initial age a_0 and live at most a_T periods. The total number of distinct individuals born in the economy is discrete and denoted by N , where a fraction μ_h of individuals born in cohort h . The initial age, a_0 , should be interpreted as the age at which individuals exit college and enter the labor force. The age of an individual i in cohort h at time t is $a_{ht} = a_0 + t - h$. Before age a_T , individuals face age-dependent mortality risk, where the survival probability at age $a + 1$ conditional on survival age a is denoted by m_a . Between

ages a_0 and $a_R - 1$, individuals are in their working life and can supply labor to earn income. At age a_R , individuals exogenously transition to retirement and cannot supply labor.

4.1.2 Preferences

In each period of working life, individuals choose consumption, c , and labor supply, ℓ . An individual i at age a has [Epstein and Zin \(1989\)](#)-[Weil \(1990\)](#) preferences over consumption and labor supply defined recursively by:

$$V_{ia} = \left[(1 - \beta) n_a \left(\frac{c_{ia}}{n_a} - \kappa \frac{\ell_{ia}^{1+\phi^{-1}}}{1 + \phi^{-1}} \right)^{1-\sigma} + \beta (m_a E_a V_{ia+1}^{1-\gamma})^{\frac{1-\sigma}{1-\gamma}} \right]^{\frac{1}{1-\sigma}}. \quad (5)$$

In (5), β is the discount factor, σ^{-1} is the intertemporal elasticity of substitution, γ is the coefficient of relative risk aversion, ϕ is the Frisch elasticity of labor supply, κ is the labor supply scaling parameter, and n_a is an equivalence scale.¹⁷ This preference specification follows [Guvenen \(2009b\)](#) and represents a recursive generalization of [Greenwood, Hercowitz, and Huffman \(1988\)](#) (GHH) preferences. These preferences eliminate wealth effects on labor supply, meaning the marginal rate of substitution between c and l is independent of changes in c .¹⁸ This assumption is consistent with empirical evidence that finds relatively limited labor supply responses to changes in wealth ([Keane 2011](#); [Cesarini, Lindqvist, Notowidigdo, and Ostling 2017](#)). I use recursive rather than time-separable preferences to independently assess the role of risk and time preferences in my normative analyses. The equivalence scale captures the evolution of household size over the life cycle, as in [Lusardi, Michaud, and Mitchell \(2017\)](#). This generates a hump-shape in consumption over the life cycle because the marginal utility of consumption increases with n_a , and calibrated values of n_a are hump-shaped.

4.1.3 Labor Income Process

During working life, the labor income of individual i at age a , y_{ia} , is equal to the product of the individuals' wage rate, w_{ia} , and labor supply, l_{ia} , where the latter is chosen endogenously. An individuals' wage rate is modeled in partial equilibrium and consists of three components:

$$\log w_{ia} = g_{ia} + \theta_{ia} + \epsilon_{ia}. \quad (6)$$

¹⁷(5) also embeds the assumption that $u_d^{1-\gamma} = 0$, where u_d is the utility upon death. This assumption is standard in life cycle models with recursive preferences. However, with $\gamma > 1$, it implies that $u_d = \infty$. [Bommier, Harenberg, Le Grand, and O'Dea \(2020\)](#) point out some undesirable implications of this assumption in models where mortality is endogenous, which is not the case in my model.

¹⁸[Auclert and Rognlie \(2017\)](#) point out that GHH preferences generate an additional source of amplification in response to shocks due to the complementarity of consumption and leisure. This amplification is not present in my model since wage rates are not determined in general equilibrium.

The first component, g_{ia} , is a deterministic life cycle component whose specific form is discussed later. The other two components, θ_a and ϵ_a , capture stochastic components of an individuals' wage process, which take the following forms:

$$\begin{aligned} \theta_{ia} &= \rho\theta_{ia-1} + \alpha \log(l_{ia-1}) + \nu_{ia}, & \theta_{ia_0} &= \delta_i, \\ \delta_i &\sim \mathcal{N}(0, \sigma_i^2), & \nu_{ia} &\sim \mathcal{N}(0, \sigma_\nu^2), & \epsilon_{ia} &\sim \mathcal{N}(0, \sigma_\epsilon^2). \end{aligned} \quad (7)$$

This specification in (7) allows for permanent and transitory shocks to wage rates, which is important because individuals can only self-insure against the latter in incomplete markets (Blundell, Pistaferri, and Preston 2008). The transitory component consists solely of ϵ_{ia} , which is i.i.d. within and across individuals. The permanent component is captured by θ_{ia} , which depends on three factors. First, it depends on permanent shocks ν_{ia} , which have persistence captured by ρ . Second, it depends on an individual fixed effect, δ_i , which captures ex-ante heterogeneity across individuals. Finally, it exhibits learning-by-doing following Keane (2016), in which past values of labor supply affect future wage rates with elasticity α . Although I find little evidence for learning-by-doing affecting bunching empirically, I estimate a version of the model with α set based on prior literature to assess its importance in policy counterfactuals.

Aside from the presence of learning-by-doing and the fact that θ_{ia} is not a random walk, this specification of the wage rate process is similar to the standard permanent-transitory income processes used in canonical life cycle models of consumption and portfolio choice (Zeldes 1989; Gourinchas and Parker 2002; Cocco, Gomes, and Maenhout 2005). A key difference, however, is that the *income* process is endogenous because individuals choose their labor supply.

4.1.4 Education Levels

In addition to having different initial permanent income through δ_i , individuals differ ex-ante based on their education levels. There are two education levels denoted by $\mathcal{E}_i \in \{0, 1\}$, where

$$\mathcal{E}_i \sim \text{Bernoulli}(p_E). \quad (8)$$

Individuals with $\mathcal{E}_i = 1$ are referred to as “Graduates”, meaning they have a college degree, while those with $\mathcal{E}_i = 0$ are referred to as “Non-Graduates”. Individuals' education level determines the deterministic component of their income process, g_{ia} , which takes the following form:

$$g_{ia} = \delta_0 + \delta_1 a + \delta_2 a^2 + \mathcal{E}_i (\delta_0^E + \delta_1^E a). \quad (9)$$

This specification captures that the returns to experience are quadratic (in logs), as in [Mincer \(1974\)](#), and that Graduates may have different wage levels and profiles.¹⁹

4.1.5 Labor Supply Optimization Frictions

Individuals choose their labor supply at the same time they choose consumption, which occurs at the end of each period after all shocks are realized. I introduce optimization frictions that prevent individuals from frictionlessly choosing their labor supply. As discussed in Section 3.8, these frictions are needed to generate individuals above the repayment threshold. Because isolating the importance of every possible optimization friction is not possible given the available data and empirical variation, I instead follow [Nakamura and Steinsson \(2010\)](#) and [Andersen, Campbell, Nielsen, and Ramadorai \(2020\)](#) and consider a specification that nests the two canonical types of imperfect adjustment: state-dependent and time-dependent adjustment.²⁰

The first optimization friction is that choosing labor supply in the current period that is different from the past period, $\ell_{ia} \neq \ell_{ia-1}$, requires paying a fixed cost of f , except in individuals' first period of life. This fixed cost generates (S, s) -type behavior and makes labor supply adjustment state-dependent, meaning individuals only adjust their labor supply when the benefits of adjustment are sufficiently high.²¹ This cost could capture real costs associated with changing labor supply, such as search costs associated with changing jobs when hours are constrained by firms, or psychological costs, such as the hassle costs of adjusting a working hours schedule. The fixed cost is modeled as a utility cost, as axiomatized by [Masatlioglu and Ok \(2005\)](#).

The second optimization friction is that only a fraction λ of individuals in each period can adjust their labor supply à la [Calvo \(1983\)](#).²² Formally, individuals with $\omega_{ia} = 1$ can choose consumption

¹⁹I do not allow for the possibility that the quadratic component of g_{ia} differs with \mathcal{E}_i . This is because *ALife* only covers 1991-2019 and does not have direct measures of education. Since I instead infer education level based on the presence of HELP debt, the oldest individual I observe in the sample with $\mathcal{E}_i = 1$ is around age 50-55. Without the final 5-10 years of working life, it is difficult to identify this additional parameter.

²⁰An alternative optimization friction is optimization errors (or noise) that could take two forms, both inconsistent with my empirical evidence. The first is anticipated errors, in which individuals know they cannot control labor supply perfectly. This, however, makes the prediction that there will be excess mass further to the left of the threshold as individuals reduce their labor supply even more to ensure they do not end up above it, which is not the case in [Figure 3](#). The second is unanticipated optimization errors, where labor supply equals individuals' choice plus an error. This predicts that the bunching will be diffuse around the repayment threshold (as in [Chetty et al. 2013](#)), while the bunching in [Figure 3](#) is sharp.

²¹The (S, s) model has been used to model several decisions in household finance, such as portfolio choice ([Abel, Eberly, and Panageas 2013](#)) and saving decisions ([Choukhmane 2021](#)), but also in many other settings, such as price-setting ([Caplin and Spulber 1987](#)) and capital investment ([Caballero and Engel 1999](#)).

²²This type of time-dependent adjustment provides a good fit to mortgage refinancing and retail investor portfolio rebalancing behavior ([Andersen et al. 2020](#); [Giglio, Maggiori, Stroebel, and Utkus 2021](#)).

and labor supply and those with $\omega_{ia} = 0$ can only adjust consumption, where:

$$\omega_{ia} \sim \begin{cases} 1, & \text{if } a = a_0, \\ \text{Bernoulli}(\lambda), & \text{else.} \end{cases} \quad (10)$$

This adjustment shock, ω_{ia} , generates time-dependent labor supply adjustment. Economically, this shock could capture frictions on the demand-side of the labor market that result in the slow arrival of opportunities to adjust labor supply or job transitions (as in [Kleven et al. 2023](#)). Alternatively, this could capture simple inattention, where $1 - \lambda$ captures the fraction of inattentive individuals.²³

4.1.6 Liquid Assets

At age a_0 , individuals are endowed with an initial stock of liquid assets, A_{ia_0} , where

$$A_{ia_0} \sim \begin{cases} 0, & \text{with probability } p_A(\mathcal{E}_i), \\ \text{Log-normal}(\mu_A(\mathcal{E}_i), \sigma_A(\mathcal{E}_i)^2), & \text{with probability } 1 - p_A(\mathcal{E}_i). \end{cases} \quad (11)$$

The dependence of this distribution on \mathcal{E}_i allows for the possibility that individuals with different education levels also have different initial liquidity. In subsequent periods, individuals liquid asset balances after consumption at age $a - 1$ are denoted by A_{ia} . Positive balances in the liquid asset pay a gross return of R . Individuals can also borrow using unsecured credit up to an age-dependent borrowing limit, \underline{A}_a . The interest rate on borrowing is $R + \tau_b$, where τ_b captures the borrowing rate wedge. Individuals' asset income, i_{ia} , is received prior to consumption at age a and is equal to:

$$i_{ia} = r(A_{ia}) * A_{ia}, \quad r(A_{ia}) = R - 1 + \tau_b * \mathbf{1}(A_{ia} < 0). \quad (12)$$

In this model, the interest rate and borrowing wedge are taken as exogenous. This is primarily done for tractability, but it is unlikely to affect the conclusions from my counterfactual analyses for two reasons. First, individuals with large student debt balances, who are most affected by the policy changes I consider, are young and hold a relatively small share of aggregate wealth. Second, simulation results show that the change in the aggregate stock of liquid assets in response to these policy changes is negligible, suggesting any change in the equilibrium interest rate would be small.

²³This is an imperfect way of modeling inattention because agents are sophisticated about their lack of inattention. However, modeling naive inattention introduces complications with individuals violating their budget constraints that are beyond the scope of this paper.

4.1.7 Student Debt

At age a_0 , individuals are also endowed with debt balances, D_{ia_0} , where

$$D_{ia_0} \sim \begin{cases} 0, & \text{if } \mathcal{E}_i = 0, \\ \text{Log-normal}(\mu_d, \sigma_d^2), & \text{if } \mathcal{E}_i = 1. \end{cases} \quad (13)$$

These initial debt balances are exogenous in my model because I focus on the trade-off between insurance and moral hazard ex-post. In subsequent periods, debt balances evolve according to:

$$D_{ia+1} = (1 + r_d)D_{ia} - d_{ia}, \quad d_{ia} = d(y_{ia}, i_{ia}, D_{ia}, a, t), \quad (14)$$

where r_d is the (net) interest rate on student debt and d_{ia} is the required student debt repayment that is determined by the repayment function, $d(\cdot)$. This repayment function depends on individuals' income, debt balance, and age. I assume any outstanding debt is discharged once individuals enter retirement at $a = a_R$ ²⁴ or upon death. When I estimate the model, this repayment function is set equal to the HELP repayment function in [Figure 2](#). In counterfactuals, I consider alternative repayment functions, such as those in the US and income-sharing agreements.

4.1.8 Government

The government earns revenue from progressive taxes on labor and asset income and student debt repayments. Total taxes on labor and asset income are denoted by $\tau_{ia} = \tau(y_{ia}, i_{ia}, t)$. Government expenditures include student loans to newborn individuals at $a = a_0$, means-tested unemployment benefits, $ui_{ia} = ui(y_{ia}, i_{ia}, A_{ia})$, and a means-tested retirement pension, $\bar{y}_R(A_{ia})$. The government also pays a net consumption floor, \underline{c}_{ia} , to ensure individuals' consumption exceeds their disutility from labor supply by \underline{c} in the event they do not adjust the latter.²⁵ For all government taxes and transfers, including debt repayments, there is no deduction for interest paid on unsecured borrowing.

4.1.9 Recursive Formulation

Individuals solve a stochastic dynamic programming problem, which can be formulated recursively. There are five continuous state variables: A_{ia} = beginning-of-period liquid assets, ℓ_{ia-1} = past labor supply, D_{ia} = student debt balance, θ_{ia} = persistence component of wage rate, and ϵ_{ia}

²⁴I do this because individuals' only source of income in retirement is capital income and the counterfactual repayment contracts I consider are contingent solely on wage income.

²⁵The combination of GHH preferences and labor supply optimization frictions implies that there will be parts of the state space where individuals cannot ensure consumption net of the disutility of labor supply is positive, which causes V_{ia} to be poorly-behaved. This consumption floor prevents that but is never received by any individuals in simulations.

= transitory component of wage rate. There are four discrete state variables: t = current year, a = age, \mathcal{E}_i = level of education, and ω_{ia} = Calvo adjustment shock. Denote the vector of these state variables for individual i at age a as \mathbf{s}_{ia} and $E_a(\cdot) = E(\cdot | \mathbf{s}_{ia+1})$ as the conditional expectation over the three shocks, ω_{ia+1} , ν_{ia+1} , and ϵ_{ia+1} . There are two controls: end-of-period liquid assets, A_{ia+1} , and labor supply, ℓ_{ia} , where consumption, c_{ia} is pinned down by the budget constraint.

Suppressing i subscripts, individuals at age $a < a_R$ that receive the adjustment shock and individuals at age $a = a_0$ solve the following problem:

$$V_a(\mathbf{s}_a) = \max_{A_{a+1}, \ell_a} \left\{ (1 - \beta) n_a \left[\frac{c_a}{n_a} - \kappa \frac{\ell_a^{1+\phi^{-1}}}{1 + \phi^{-1}} - f * \mathbf{1}(\ell_a \neq \ell_{a-1}) \right]^{1-\sigma} + \beta [m_a E_a (V_{a+1}(\mathbf{s}_{a+1})^{1-\gamma})]^{\frac{1-\sigma}{1-\gamma}} \right\}^{\frac{1}{1-\sigma}}$$

subject to: (6), (7), (9), (10), (12), (14), and

$$c_a + A_{a+1} = y_a + A_a + i_a - d_a - \tau_a + w_i \ell_a$$

constraints: $A_{a+1} \geq \underline{A}_{a+1}$ and $\ell_a \geq 0$

boundary conditions: (7), (8), (11), (13), and $\ell_{a_0-1} = \ell_{a_0}$

Individuals at age $a < a_R$ that do not receive the adjustment shock solve the following problem:

$$V_a(\mathbf{s}_a) = \max_{A_{a+1}} \left\{ (1 - \beta) n_a \left[\frac{c_a}{n_a} - \kappa \frac{\ell_{a-1}^{1+\phi^{-1}}}{1 + \phi^{-1}} \right]^{1-\sigma} + \beta [m_a E_a (V_{a+1}(\mathbf{s}_{a+1})^{1-\gamma})]^{\frac{1-\sigma}{1-\gamma}} \right\}^{\frac{1}{1-\sigma}}$$

subject to: (6), (7), (9), (10), (12), (14), and

$$c_a + A_{a+1} = y_a + A_a + i_a - d_a - \tau_a + w_i \ell_a + \underline{c}_a$$

constraint: $A_{a+1} \geq \underline{A}_{a+1}$

Retired individuals at age $a \geq a_R$ solve the following problem:

$$V_a(\mathbf{s}_a) = \max_{A_{a+1}} \left\{ (1 - \beta) n_a \left(\frac{c_a}{n_a} \right)^{1-\sigma} + \beta [m_a E_a (V_{a+1}(\mathbf{s}_{a+1})^{1-\gamma})]^{\frac{1-\sigma}{1-\gamma}} \right\}^{\frac{1}{1-\sigma}}$$

subject to: (12), (14), and $c_a + A_{a+1} = \bar{y}_R(A_{ia}) + A_a + i_a - \tau(0, i_a, t)$

constraint: $A_{a+1} \geq \underline{A}_{a+1}$

boundary condition: $V_{a_T}(\mathbf{s}) = (1 - \beta)^{\frac{1}{1-\sigma}} c_{a_T} \forall \mathbf{s}$

The model is solved using standard numerical discrete-time dynamic programming techniques. The code to solve and simulate the model is compiled in Intel Fortran 2018 and executed in parallel using both MPI and OpenMP across 1,536 CPU threads. For additional details on the solution technique, see Appendix F.

4.2 Estimation Procedure

This section describes how I estimate the model. In a first step, I calibrate parameters that can be observed directly and others based on prior literature. In a second step, I estimate the key parameters in the model, which are the labor supply preference parameters, discount factor, and parameters controlling the wage process, using simulated minimum distance.

4.2.1 Calibrated Parameters

Table 2 shows the values of parameters that I can calibrate directly using either observed Australian data or formulas from the Australian tax and transfer system. In what follows, I provide a brief description of this calibration; see Appendix G for additional details.

Demographics. Individuals are born at age 22, which corresponds to the typical age at which students graduate from university in Australia, retire at age 65, which is the age at which the Australian retirement pension began to be paid in 2004, and die with certainty after age 89. Prior to age 89, individuals' mortality risk is calibrated to match Australia's life tables. Cohort-specific birth rates are calibrated to match the fraction of 22-year-olds in each year in *ALife*. I use data on household sizes from HILDA to compute equivalence scales using the same procedure in [Lusardi et al. \(2017\)](#).

Interest rates and borrowing. There is no inflation in the model, and the numeraire is equal to \$1 AUD in 2005. When compared with model moments, all empirical moments are deflated to 2005 AUD using the indexation rates for HELP thresholds. The real interest rate is set to 1.84%, the (geometric) average real interest rate paid on deposits between 1991 and 2019 in Australia. The unsecured borrowing rate is set based on average credit card borrowing rates. Age-specific borrowing limits are set based on credit card limits reported in HILDA. The real interest rate on student debt is set to zero, since HELP debt has a nominal interest rate equal to inflation.

Initial conditions. The distribution of initial assets is calibrated to match the liquid wealth distribution of individuals between ages 18 and 22. The fraction of individuals with college degrees, p_E , is equal to the fraction of 22-year-old individuals in *ALife* that have positive debt balances, which is the year by which most individuals have started their undergraduate degrees in Australia. The distribution of initial debt balances is set based on the distribution of debt balances among individuals younger than age 26 in *ALife*, the age by which most individuals have finished undergraduate studies in Australia and debt balances reach their maximum in real terms.

Government taxes and transfers. Income and capital taxes are set to match the individual income tax schedules provided by the ATO in 2004 and 2005. Unemployment benefits are means-tested and calculated based on the Newstart Allowance, the primary form of government-provided

income support in Australia to individuals above 22. The retirement pension is calculated following the Age Pension formula, the primary government-provided form of income-support to retirees in Australia. The age pension is available to individuals at age 65 and is means-tested based on assets and income.

Preference parameters. The preference parameters I do not estimate due to a lack of identifying variation are relative risk aversion and the elasticity of intertemporal substitution. I set $\gamma = \sigma = 2.23$ based on [Choukhmane and de Silva \(2023\)](#), which corresponds to time-separable preferences with a relative risk aversion of 2.23 and an EIS of $2.23^{-1} = 0.45$. In counterfactuals, I consider the effects of increasing γ and decreasing σ to the calibration used in [Bansal and Yaron \(2004\)](#), which introduces a preference for early resolution of uncertainty.

Learning-by-doing. My data also do not provide sufficient variation to identify the learning-by-doing parameter, α , which captures the elasticity of future wages to current labor supply. This is because learning-by-doing has a minimal effect on individuals' incentives to bunch below the repayment threshold due to the envelope theorem. I thus consider two different values of $\alpha \in \{0, 0.24\}$, where the latter corresponds to the median value from the meta-analysis conducted by [Best and Kleven \(2012\)](#). I consider $\alpha = 0$ as my baseline model and only discuss results from the latter, where they differ significantly from the former.

4.2.2 Simulated Minimum Distance Estimation

I estimate the remaining 14 parameters that cannot be calibrated directly using simulated minimum distance, which I denote by Θ :

$$\Theta = \left(\underbrace{\phi \quad f \quad \lambda \quad \kappa \quad \beta}_{\text{preference parameters}} \quad \underbrace{\delta_0 \quad \delta_1 \quad \delta_2 \quad \delta_0^E \quad \delta_1^E}_{\text{wage profile parameters}} \quad \underbrace{\rho \quad \sigma_\nu \quad \sigma_\epsilon \quad \sigma_i}_{\text{wage risk parameters}} \right).$$

These parameters can be divided into three groups: preference parameters, parameters governing the age profile of wages, g_{ia} , and finally, parameters governing shocks to the wage process. In contrast to the standard approach of estimating life cycle models of consumption and portfolio choice (e.g., [Gourinchas and Parker 2002](#); [Cocco et al. 2005](#)), I cannot estimate the latter two sets of parameters separately in a first stage because my income process is endogenous. I thus proceed by combining a standard set of moments used to identify the latter two sets of parameters in models with exogenous income processes with the quasi-experimental variation from the policy change to the HELP repayment function. As detailed in the Section 4.2.3, individuals' responses to this policy change are what allow me to separately identify the three most important parameters: the labor supply elasticity, ϕ , fixed adjustment cost, f , and Calvo adjustment probability, λ .

Table 2. Values of Calibrated Model Parameters

Description	Parameter(s)	Values/Targets
Demographics		
Age in first year of life	a_0	22
Age in first year of retirement	a_R	65
Age in final year of life	a_T	89
Mortality rates	$\{m_a\}$	APA Life Tables
First and last cohorts	\underline{h}, \bar{h}	1963, 2019
Cohort birth probabilities	$\{\mu_h\}$	ALife
Equivalence scale	$\{n_a\}$	HILDA Household Size
Number of distinct individuals	N	1,600,000
Year of simulated policy change	T^*	2005
Assets		
Numeraire	\cdot	\$1 AUD in 2005
Real interest rate	$R - 1$	1.84%
Unsecured borrowing wedge	τ_b	14.6%
Borrowing limit	$\{\underline{A}_a\}$	HILDA Credit Card Limit
Probabilities of zero initial assets	$p_A(1), p_A(0)$	0.197, 0.350
Means of $\log A_{ia_0}$	$\mu_A(1), \mu_A(0)$	7.42, 6.79
Standard deviations of $\log A_{ia_0}$	$\sigma_A(1), \sigma_A(0)$	1.72, 2.64
Preference Parameters		
Relative risk aversion	γ	2.23
Elasticity of intertemporal substitution	σ^{-1}	0.45
Learning-by-doing parameter	α	0, 0.24
Student Debt		
Fraction of Graduates	p_E	0.308
Real interest rate on debt balances	r_d	0%
Means of $\log D_{ia_0}$	μ_d	9.40
Standard deviations of $\log D_{ia_0}$	σ_d	0.86
Debt repayment function	$d(\cdot)$	HELP 2004 at $t < T^*$, HELP 2005 at $t \geq T^*$
Government		
Income and capital taxes	$\tau(\cdot)$	ATO Income Tax Formulas
Unemployment benefits	$ui(\cdot)$	ATO Newstart Allowance
Retirement pension	$\bar{y}_R(\cdot)$	ATO Age Pension
Net consumption floor	\underline{c}	\$40

Simulated policy change. I replicate the policy change shown in [Figure 2](#) within the model by solving the model for two different specifications of the student debt repayment function, $d(\cdot)$: (i) the HELP 2004 repayment formula and (ii) the HELP 2005 repayment formula. Starting at $t = \underline{h} = 1963$, I simulate cohorts of individuals making choices under the 2004 formula. At $t = T^* = 2005$, I then conduct a one-time unanticipated policy change in which all existing debtholders born at $t < T^*$ and subsequent debtholders start repaying under the 2005 formula.

Estimator. I estimate the vector of parameters, Θ , using simulated minimum distance. This procedure consists of choosing a set of estimation targets, which is a vector of moments, summary statistics, or auxiliary parameters, and a weighting matrix. Denote the empirical values of estimation targets as \hat{m} , the vector of the estimation targets estimated in the model via simulation at parameters Θ as $m(\Theta)$, and the weighting matrix as $W(\Theta)$. My estimate of Θ is then defined as Θ^* , where

$$\Theta^* = \arg \min_{\Theta} (\hat{m} - m(\Theta))' W(\Theta) (\hat{m} - m(\Theta)).$$

I choose $W(\Theta)$ so that this objective function equals the sum of squared arc-sin deviations between \hat{m} and $m(\Theta)$. The 47 estimation targets I use are shown in [Table 3](#) and discussed in the next section. Additional details on the calculation of each target in the data and within the model are presented in [Appendix H](#). Because this optimization problem is high-dimensional and likely has several local optima, I perform the minimization using a modified version of the TikTak algorithm introduced by [Arnoud, Guvenen, and Kleineberg \(2019\)](#) detailed in [Appendix H](#).

4.2.3 Selection of Estimation Targets and Parameter Identification

I now discuss how each parameter is identified by the estimation targets in my simulated minimum distance estimation. All parameters are of course, jointly identified, but I choose the set of estimation targets so that each one is most sensitive to a subset of parameters. [Table 3](#) lists each estimation target and the parameter(s) that it primarily identifies. The discussion in this section is mostly qualitative; [Table A2](#) provides the elasticities of each estimation target with respect to each structural parameter as supporting evidence.

Labor supply elasticity, ϕ . The labor supply elasticity is identified by the extent of bunching in the HELP Income distribution below the repayment thresholds both before and after the policy change: a larger elasticity implies greater mass below these thresholds. To characterize this bunching, I use the real distributions of HELP Income among debtholders three years before and three years after the policy change. I focus on the distribution within \$3,000 of the repayment thresholds so that these targets are primarily affected by the labor supply elasticity rather than wage profile parameters. I use bins of \$500 to minimize approximation error in my estimation of the model moments.

Fixed adjustment cost, f , and Calvo probability, λ . These two optimization frictions are jointly identified by the number of individuals above the repayment threshold: even with a very small labor supply elasticity, a model $f = 0$ and $\lambda = 1$ would predict no individuals immediately above the repayment threshold because locating below it increases cash-on-hand. To separately identify these two parameters, I exploit the fact that adjustment costs imply *state*-dependent labor supply responses. In particular, adjustment costs predict disproportionately more bunching at the 2005 repayment threshold relative to the lowest 2005 0.5% threshold because the former has a discontinuity in repayment rate of 4% rather than 0.5%. Additionally, adjustment costs generate larger bunching among individuals with more debt, for whom the present discount value of reducing labor supply is larger. In contrast, a model with pure Calvo adjustment implies less heterogeneity in bunching across debt balances because adjustment depends on whether individuals receive the adjustment shock.

To characterize bunching at different thresholds and among individuals with different debt balances using a manageable number of estimation targets, I compute the ratio of individuals within \$250 below to within \$250 above each threshold in each sample. This ratio captures the extent of bunching: more bunching implies more individuals below relative to above the threshold and thus a higher ratio. To target heterogeneity across thresholds with different repayment rates, I compute this ratio at the 2004 threshold prior to the policy change, at the 2005 threshold after the policy change, and at the lowest 2005 0.5% threshold after the policy change (see [Figure A17](#) for a comparison of the latter two). I then compute it at the 2005 threshold after the policy change among individuals in the bottom and top quartile of debt balances (within each year) to target heterogeneity in responses across debt balances.

Labor supply scaling parameter, κ . This parameter is simply a scaling parameter that determines the scale of l_{ia} . It is identified by the average value of l_{ia} , which I normalize to one. A higher value increases the disutility for labor supply and thus lowers average values of l_{ia} .

Time discount factor, β . The time discount factor is identified by the average level of capital income. A higher value makes individuals more patient, increasing saving and hence capital income. I target capital income between ages 40 and 44, the midpoint of individuals' working lives.

Wage profile parameters, δ_0 , δ_1 , δ_2 , δ_0^E , and δ_1^E . These parameters are primarily identified by the regressions of log income onto polynomials in age and the education level dummy, in addition to the average level of income. If labor supply was exogenous, they could be estimated separately using these moments alone. However, with endogenous labor supply, these parameters control the *wage* rather than the income process and must be estimated jointly because the former is not observable.

Wage risk parameters, ρ , σ_ν , σ_ϵ , and σ_i . These parameters are identified by how the cross-sectional variance of log income varies with age and the percentiles of income growth at one-

year and five-year horizons. This set of moments is standard in the literature used to estimate exogenous income processes (e.g., [Guvenen, McKay, and Ryan 2022](#)), and the identification works similarly here, even though the income process is endogenous. The cross-sectional variance at age 22 identifies σ_i , the variance of the initial permanent income. The extent to which the cross-sectional variance increases with age identifies the persistence of income shocks, ρ : more persistent shocks generate a greater increase in variance over the life cycle ([Deaton and Paxson 1994](#)). The sum of the variances of permanent and transitory income shocks, σ_ν and σ_ϵ , are identified by the level of this cross-sectional variance at later ages. These two variances are then separated using the percentiles of future income growth: a larger variance of permanent shocks, σ_ν , delivers fatter tails in 5-year relative to 1-year income growth.

Table 3. List of Estimation Targets in Simulated Minimum Distance Estimation

Estimation Target	Parameter(s) Most Sensitive to Target
Labor Supply Preference Parameters	
Real distribution of HELP Income among debtholders in 2002-2004 within \$3000 of the 2004 repayment threshold in bins of \$500	ϕ, f, λ
Real distribution of HELP Income among debtholders in 2005-2007 within \$3000 of the 2005 repayment threshold in bins of \$500	ϕ, f, λ
Ratio of number of individuals with HELP Income within \$250 below to the number within \$250 above the 2004 repayment threshold in 1998-2004	ϕ, f, λ
Ratio of number of individuals with HELP Income within \$250 below to the number within \$250 above the 2005 repayment threshold in 2005-2018	ϕ, f, λ
Ratio of number of individuals with HELP Income within \$250 below to the number within \$250 above the 2005 repayment threshold in 2005-2018 among individuals in the bottom and top quartile of debt balances in each year	f, λ
Ratio of number of individuals with HELP Income within \$250 below to the number within \$250 above the lowest 2005 0.5% threshold in 2005-2018	f, λ
Other Preference Parameters	
Average labor supply of employed individuals	κ
Average capital income between ages 40 and 44	β
Wage Profile Parameters	
Average salary & wages of employed individuals	δ_0
Regression coefficients of log salary & wages onto a and a^2	δ_1, δ_2
Regression coefficients of log salary & wages onto \mathcal{E}_i and $\mathcal{E}_i a$ among $h \geq 1991$	δ_0^E, δ_1^E
Wage Risk Parameters	
Within-cohort cross-sectional variance of log salary & wages at age 22	σ_i
Within-cohort cross-sectional variance of log salary & wages at ages 32, 42, 52, and 62	$\rho, \sigma_\nu, \sigma_\epsilon$
10th and 90th percentiles of 1-year and 5-year salary & wages growth	$\sigma_\nu, \sigma_\epsilon$

Notes: Additional details are presented in [Appendix H](#).

4.3 Baseline Estimation Results and Model Fit

The results for my baseline simulated minimum distance estimation are reported in column (1) of [Table 4](#). My baseline estimate of the (Frisch) labor supply elasticity is 0.114. This estimate is

close to the mean value of 0.15 reported for Hicksian intensive-margin labor supply elasticities, which corresponds to ϕ in my model given the absence of wealth effects, from a meta-analysis of hours and taxable income responses in [Chetty \(2012\)](#). The baseline estimation also delivers a fixed adjustment cost of \$377, which corresponds to around 1% of average income in the model, and Calvo adjustment probability of 0.183. This value of the Calvo parameter implies that, in expectation, individuals receive an opportunity to adjust their labor supply every 5.4 years.


Table 4. Simulated Minimum Distance Estimation Results

Parameter		Estimation				
		(1)	(2)	(3)	(4)	(5)
Labor supply elasticity	ϕ	0.114 (.004)	0.005 (.000)	0.188 (.003)	0.053 (.002)	0.082 (.002)
Fixed adjustment cost	f	\$377 (\$13)	\$0 .	\$2278 (\$21)	\$0 .	\$762 (\$10)
Calvo parameter	λ	0.183 (.003)	1 .	1 .	0.147 (.002)	0.346 (.009)
Labor supply scaling parameter	κ	0.560 (.007)	0.030 (.003)	0.059 (.014)	0.510 (.012)	1.242 (.116)
Time discount factor	β	0.973 (.001)	0.996 (.000)	0.972 (.001)	0.944 (.001)	0.951 (.001)
Wage profile parameters	δ_0	8.922 (.009)	9.862 (.002)	8.680 (.006)	9.389 (.007)	9.197 (.007)
	δ_1	0.073 (.000)	0.111 (.000)	0.073 (.000)	0.063 (.000)	0.070 (.000)
	δ_2	-0.001 (.000)	-0.002 (.000)	-0.001 (.000)	-0.001 (.000)	-0.001 (.000)
	δ_0^E	-0.487 (.002)	-0.294 (.000)	-0.450 (.001)	-0.530 (.002)	-0.480 (.002)
	δ_1^E	0.020 (.000)	0.032 (.000)	0.018 (.000)	0.021 (.000)	0.018 (.000)
	Persistence of permanent shock	ρ	0.930 (.000)	0.914 (.000)	0.943 (.000)	0.922 (.000)
	Standard deviation of permanent shock	σ_ν	0.236 (.000)	0.076 (.000)	0.196 (.000)	0.268 (.000)
Standard deviation of transitory shock	σ_ϵ	0.130 (.000)	0.504 (.000)	0.168 (.000)	0.077 (.002)	0.064 (.002)
Standard deviation of individual FE	σ_i	0.599 (.003)	0.101 (.001)	0.541 (.003)	0.654 (.003)	0.625 (.003)
Learning-by-doing parameter		α	0	0	0	0.24

Notes: This table shows the results from simulated minimum distance estimations, where each column corresponds to a separate estimation. Entries in the table correspond to parameter estimations with standard errors presented below in parentheses. All estimations use the same set of estimation targets in [Table 3](#). Parameters that are fixed at their respective values and not estimated are indicated with “.” in place of a standard error.

[Figure 9](#) shows how the baseline model fits the distribution of HELP Income before and after the policy change. The model provides a close approximation of both distributions, especially the mass of individuals immediately below and above the two repayment thresholds. There are slight differences at other points that reflect the fact that the model cannot perfectly match the age profile of income.

Figure 9. Baseline Model Fit: HELP Income Distribution around Policy Change




Notes: Bars represent 95% confidence intervals based on bootstrapped standard errors with 1000 iterations.

Figure 10 illustrates the model's fit of the amount of bunching at other repayment thresholds, in addition to among individuals with different debt balances. Consistent with Figure 9, the model replicates the bunching at the 2004 and 2005 repayment thresholds well. However, the model can also replicate the relatively small amount of bunching at the lowest 0.5% repayment threshold after the policy change. The presence of a fixed adjustment cost is crucial for this result: in a model with only Calvo adjustment, there is less of a difference in the amount of bunching at this threshold and the 0% threshold because the probability that individuals receive an adjustment shock is independent of their level of income. Similarly, the presence of an adjustment cost helps match the difference in bunching between individuals with low and high debt balances. Quantitatively, the model misses slightly on matching the right amount of bunching at the 0.5% threshold. This is because matching this moment better would require performing worse on the others: increasing the adjustment cost would improve the fit at the lower 0.5% threshold but would also further decrease the amount of bunching among individuals with low debt balances, which the model already underestimates.

Table 5 shows the fit of the model on the remaining target moments, which are primarily used to estimate the remaining parameters outside of the labor supply elasticity, fixed adjustment cost, and Calvo parameter. The model provides a relatively good fit to average labor income and the age profiles of income, which are most affected by the wage profile parameters in Table 4. The fit is not perfect because income in the model depends on endogenous labor supply. To the extent the age profile of labor supply varies over the life cycle for reasons outside the model, it will be unable to match these age profiles.

Table 5 shows the cross-sectional variance of income is mostly increasing over the life cycle, a fact

Figure 10. Baseline Model Fit: Bunching around Thresholds



Notes: Bars represent 95% confidence intervals based on bootstrapped standard errors with 1000 iterations.

first documented by [Deaton and Paxson \(1994\)](#). The model can replicate this pattern due to the high persistence of permanent shocks, $\rho = 0.93$. [Guvenen \(2009a\)](#) points out that such an estimate is upward-biased in models without heterogeneous income profiles. My model features a limited form of profile heterogeneity across the two education groups, which brings my estimate of ρ down below typical unit root estimates in models with homogenous income profiles ([Guvenen 2009a](#)). Nevertheless, to address the concern that an upward bias in ρ would overstate the income risk and hence the insurance benefits from income-contingent loans, I consider alternative values of ρ when comparing different repayment policies.

Finally, the model generates a similar level of capital income for individuals in middle age to the data. This moment primarily identifies the annual discount factor, β , estimated at 0.973 in the baseline. This estimate is similar to typical estimates in life cycle models that target consumption data explicitly (e.g., [Gourinchas and Parker 2002](#)) and is less than R^{-1} . The latter finding implies individuals face a trade-off between wanting to consume at young ages due to impatience and accumulating precautionary savings, which generates buffer-stock behavior ([Carroll and Kimball 1996](#); [Carroll 1997](#)).

4.4 Identification of Labor Supply Elasticity and Optimization Frictions

The three most important parameters in my model are the labor supply elasticity, ϕ , Calvo adjustment probability, λ , and fixed adjustment cost f . [Figure A18](#) plots the simulated minimum distance objective function across these three parameters, which exhibits a clear (local) minimum. This illustrates that my estimated targets discussed in [Section 4.2.3](#) provide enough variation to separate

Table 5. Baseline Model Fit: Other Estimation Targets

Estimation Target	Data	Model
Average Labor Income	42639.373	45581.953
Cross-Sectional Variance of Log Labor Income at Age 22	0.453	0.462
Cross-Sectional Variance of Log Labor Income at Age 32	0.555	0.491
Cross-Sectional Variance of Log Labor Income at Age 42	0.577	0.525
Cross-Sectional Variance of Log Labor Income at Age 52	0.539	0.580
Cross-Sectional Variance of Log Labor Income at Age 62	0.608	0.657
Linear Age Profile Term	0.077	0.080
Quadratic Age Profile Term	-0.001	-0.001
Education Income Premium Constant	-0.574	-0.554
Education Income Premium Slope	0.023	0.023
10th Percentile of 1-Year Labor Income Growth	-0.387	-0.392
10th Percentile of 5-Year Labor Income Growth	-0.667	-0.705
90th Percentile of 1-Year Labor Income Growth	0.415	0.393
90th Percentile of 5-Year Labor Income Growth	0.698	0.710
Average Labor Supply	1.000	0.963
Average Capital Income between Ages 40 and 44	1338.846	1332.459

these different parameters that jointly determine individuals' labor supply responses. Additionally, [Figure A18](#) shows that the objective function is very smooth, which lends confidence to my numerical solution technique. A large number of simulations (1.6 million individuals over 68 years) and the fact that no choice variables are discretized in the solution (discussed in [Appendix F](#)) are both important for generating this smoothness.

To illustrate the importance of each optimization friction, I estimate three additional models. Column (2) of [Table 4](#) and [Figure A19](#) show the estimation results and fit of a model with no frictions (i.e., $f = 0$ and $\lambda = 1$). This estimation delivers an unreasonably low estimate of the labor supply elasticity, $\phi = 0.005$, and cannot fit most of the moments in the data. Column (3) and [Figure A20](#) show the results for a model with only a fixed adjustment cost (i.e., $\lambda = 1$). This model delivers a more reasonable estimate of the labor supply elasticity but overpredicts the amount of bunching after the policy change. This is because the fixed adjustment cost that rationalizes the amount of bunching at other thresholds is too small to prevent more individuals from bunching at the 2005 repayment threshold, which has the largest change in repayment rate.

Finally, column (4) and [Figure A21](#) show the results from a model with no fixed adjustment cost (i.e., $f = 0$). These estimation results are the closest of the three additional models to the baseline model in column (1), but this model struggles to match two key features of the data. First, the model generates too much bunching at the 0.5% threshold, which pushes the estimation to a lower value of ϕ . The intuition for this is straightforward: without a fixed adjustment cost, labor supply adjustment depends on whether an individual receives the Calvo shock, which is equally likely around all repayment thresholds. The small fixed adjustment cost in column (1) helps reduce the amount of bunching at the 0.5% threshold because the cost outweighs the benefit for many individuals while still being too small to affect the bunching at other thresholds where the benefit

is larger. In order to compensate for the lower ϕ , which in turn predicts too little bunching at other thresholds, the estimation delivers a lower β to increase the amount of bunching. However, this lower estimate of the discount factor then causes the model to miss on a second key moment: it underestimates the amount of wealth accumulation.

4.5 Additional Discussion of Estimation Results

Model-implied Laffer curve. One way to compare my model with traditional models of labor supply is to compute the Laffer curve. In a static frictionless model of labor supply, the revenue-maximizing linear tax rate is $\frac{1}{1+\phi}$ (Saez 2001), which is 90% given my estimate of $\phi = 0.11$. [Figure A23](#) plots the Laffer curve in my model and shows the revenue-maximizing tax rate is around 80%. This suggests that my model delivers reasonable estimates for the effects of income taxation despite the fact that it was estimated to match the effects of income-contingent loans.

Wage risk parameters. A growing literature uses administrative data to estimate parametric models of labor income risk (see e.g., Guvenen, Ozkan, and Song 2014; Guvenen, Karahan, Ozkan, and Song 2021; Braxton, Herkenhoff, Rothbaum, and Schmidt 2021; Guvenen et al. 2022; Catherine 2022). These income processes generally contain a richer set of stochastic shocks than individuals face in my model, which I abstract from due to computational constraints that arise with an endogenous income process. Nevertheless, it is instructive to compare my parameter estimates with those in the baseline specification from Guvenen et al. (2022), who estimate a similar exogenous income model using US data.

My estimate of the standard deviation of the individual fixed effect is 0.60, which is lower than 0.77 in Guvenen et al. (2022). This primarily reflects that the cross-sectional standard deviation of income at age 22 in Australian data is around 20% lower than in US data. Additionally, I estimate a standard deviation of transitory shocks that is around 30% smaller, which reflects the combination of two forces. First, the cross-sectional variance of income is lower, and the 10th/90th percentiles of income growth are less dispersed in Australia than in the US. Second, the fact that labor supply is endogenous implies that some transitory variation in income arises endogenously from labor supply adjustments rather than exogenous transitory wage shocks.²⁶ Lastly, my estimate of the standard deviation of permanent shocks is around three times as large. In addition to differences in data, this primarily reflects that I estimate $\rho = 0.93$ rather than imposing $\rho = 1$. This lower value of ρ partly reflects that I have heterogeneous income profiles across the two education groups (Guvenen 2009a), which requires a larger variance of permanent shocks to match the percentiles of 5-year income growth.

Model with learning-by-doing. Column (5) of [Table 4](#) and [Figure A22](#) show the results from

²⁶The fact that labor supply endogenously creates more volatility in income reflects the fact that my preferences have no wealth effects. In my baseline model, the ratio of the pooled variance of wage rates to income is 77%.

estimating a model with the learning-by-doing parameter, α , set equal to the median value from the meta-analysis in [Best and Kleven \(2012\)](#). The results show that this model fits the data worse than the baseline model, in particular on the heterogeneity in bunching by debt balances and the average levels of labor and capital income. The estimation of this model delivers a relatively similar labor supply elasticity but a higher estimate of the adjustment cost and Calvo parameter. This is because learning-by-doing makes bunching significantly costly for younger rather than older people: the reduction in human capital is less important for older borrowers who have fewer periods to benefit from it. As a result, this model predicts the amount of bunching increases with age, in contrast to the data (and baseline model). Additionally, since debt balances are negatively correlated with age, this model predicts too little heterogeneity in bunching with debt balances at the values of f and λ in column (1). Therefore, the estimation increases f and λ to make adjustment more state-dependent, increasing heterogeneity in bunching with debt balances.

5 Analysis of Model Mechanisms

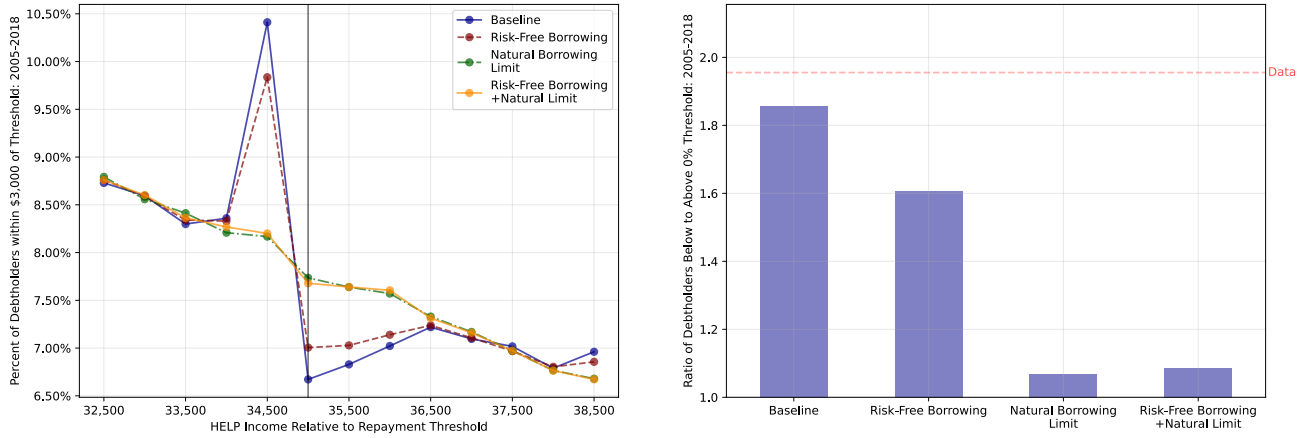
5.1 Borrowing Constraints Amplify Labor Supply Responses

[Figure 7](#) shows individuals with less wealth in the form of retirement savings are more likely to bunch below the repayment threshold, which suggests an important role for liquidity in driving labor supply responses. [Figure A24](#) shows a similar pattern holds in my estimated model: the amount of bunching decreases monotonically in individuals' initial assets. In the model, this is because additional wealth diminishes the importance of borrowing constraints by providing resources to smooth income shocks and reducing precautionary saving.

I now use the estimated model to quantify the importance of borrowing constraints directly. [Figure 11](#) shows how the amount of bunching below the repayment threshold in the model varies depending on the degree of market incompleteness individuals face. The left and right panel plot the income distribution and the ratio of individuals below to above the 2005 repayment threshold, respectively, for the baseline model, and three counterfactuals. The first counterfactual, Risk-Free Borrowing, eliminates the extra interest paid on borrowing by setting $\tau_b = 0$. Comparing this result with the baseline, the amount of bunching decreases moderately: the amount of individuals below relative to above the threshold decreases from 1.85 to 1.6, where 1 corresponds to no bunching. The second counterfactual, Natural Borrowing Limit, relaxes individuals' borrowing constraints, $\{\underline{A}_a\}$, to the natural borrowing limit.²⁷ In this counterfactual, the amount of bunching is reduced almost entirely. The third counterfactual, Risk-Free Borrowing + Natural Limit, shows that additionally setting $\tau_b = 0$ at the natural borrowing limit delivers similar results.

²⁷The natural borrowing limit cannot be computed analytically in my model. I approximate it numerically and find it corresponds to relaxing the baseline borrowing constraint by around a factor of four.

Figure 11. The Role of Borrowing Constraints in Estimated Model



Notes: The right panel plots the number of debtholders within \$250 below divided by the fraction of debtholders within \$250 above the repayment threshold over years 2005-2018.

In sum, the model predicts that labor supply responses to income-contingent repayment would be almost non-existent without borrowing constraints. This is consistent with the importance of incomplete markets in driving responses to other social insurance programs, such as unemployment insurance (Chetty 2008), mortgage default (Ganong and Noel 2022), and consumer bankruptcy (Indarte 2023), and cautions against estimating models of labor supply without frictions in credit or insurance markets.

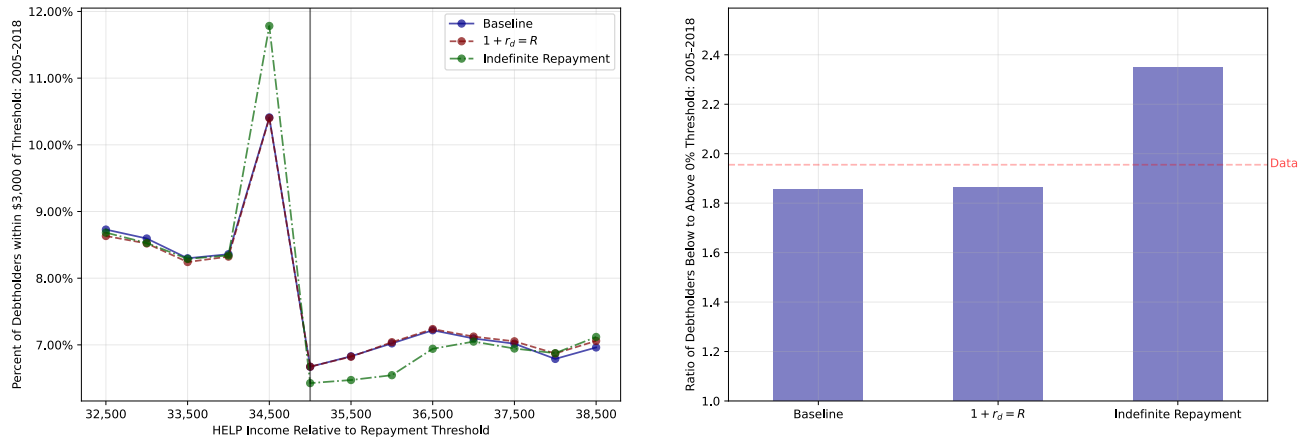
5.2 Income-Contingent Loans Generate Different Responses than Taxes

I next use my model to study how labor supply responses to income-contingent loans differ from income taxes. An extensive literature has characterized the mechanisms through which labor supply responds to taxes (Keane 2011; Saez, Slemrod, and Giertz 2012). However, there are two key differences between an income-contingent loan and an income tax in a dynamic model. First, when the interest rate on the income-contingent loan is lower than the interest rate on borrowing, income-contingent loans provide an additional incentive to reduce labor supply because doing so effectively allows individuals to borrow at a lower rate. Second, payments on income-contingent loans are capped based on individuals' initial debt balances. This reduces the incentive to reduce labor supply because the reduction in payments from doing so partly reflects a transfer over time rather than a permanent reduction.

Figure 12 quantifies the importance of these two channels by plotting the bunching around the repayment threshold in two counterfactuals. The first counterfactual eliminates the interest rate wedge between debt and liquid assets by setting $1 + r_d = R$. Comparing these results with the baseline shows that the amount of bunching is basically unchanged. The second counterfactual

eliminates the constraint that repayments are set to zero after an individuals' debt balance is paid off, but uses the same HELP repayment formula. This effectively makes income-contingent loan repayments an income tax or equity contract, where payments continue indefinitely. This counterfactual has a large effect on labor supply responses, generating almost twice as bunching below the repayment threshold. Empirically, this result is consistent with the fact that bunching increases in debt balances (Figure 6) and highlights an important way in which the effects of income-contingent loans on labor supply differ from those of income taxes.

Figure 12. Effect of Differences between Income-Contingent Loan and Tax in Estimated Model



Notes: The right panel plots the number of debtholders within \$250 below divided by the fraction of debtholders within \$250 above the repayment threshold over years 2005-2018.

6 Normative Analysis of Repayment Contracts

This section uses my estimated structural model normatively to study the welfare and fiscal impact of alternative repayment contracts. In this analysis, I take the perspective of a social planner who maximizes borrowers' expected lifetime utility by choosing one mandatory repayment contract, holding fixed borrowing choices and all prices (e.g., wages and interest rates). This problem of choosing a single repayment contract is faced by governments that only offer one contract, such as Australia and the UK. Additionally, this choice reflects the fact that my model does not capture endogenous contract selection across borrowers. In my baseline analysis, I focus on subsidized repayment contracts with a zero interest rate, as is the case in Australia.²⁸

My analysis proceeds in two steps. First, I compare *existing* income-contingent with fixed repayment contracts, which is not a budget-neutral comparison. Second, I compute the welfare gains from constrained-optimal income-contingent contracts with the same fiscal cost as fixed repayment contracts.

²⁸Under the new income-driven repayment plan in the US, known as SAVE, loan balances do not grow for individuals making their required monthly payments. Therefore, the interest rate is effectively zero for many borrowers.

6.1 Existing Income-Contingent Loans Improve Welfare, but Have Higher Fiscal Costs than Fixed Repayment

I begin by computing the welfare and fiscal impacts of various repayment contracts used in Australia and the US relative to a 25-year fixed repayment contract. This fixed repayment contract corresponds to a standard debt contract in which individuals make constant repayments over the 25 years post-graduation to repay their loan principal. I choose this contract as the benchmark because it is available in the US and has a similar duration to income-contingent contracts but differs in that it is not income-contingent.

Definition of government budget. I define the government budget, \mathcal{G} , as the expected discounted value of debt repayments and taxes net of government transfers and new debt issuance over individuals' lifetimes.²⁹ Formally,

$$\mathcal{G} \equiv \mathbf{E}_0 \left(\sum_{a=a_0}^{a_T} \underbrace{\frac{\tau_{ia} - u_{ia} - c_{ia}}{\mathcal{R}_a}}_{\text{taxes and transfers}} + \underbrace{\frac{d_{ia}}{\mathcal{R}_a} - D_{ia_0}}_{\text{debt repayments}} \right), \quad (15)$$

where $\mathbf{E}_0(\cdot)$ denotes an expectation taken over all states, including the initial state.³⁰ \mathcal{R}_a denotes the government discount rate of payments made at age a relative to a_0 , which I set equal to:

$$\mathcal{R}_a = \beta^{-(a-a_0)} \prod_{s=0}^{a-a_0} m_s. \quad (16)$$

I set \mathcal{R}_a equal to individuals' discount rates between a_0 and a , including discounting due to time preferences and mortality risk, for two reasons. First, a choice of \mathcal{R}_a different from individuals' time preferences allows the government to increase welfare simply by shifting around deterministic payments over time to take advantage of differences in discount rates. Because my analysis focuses on comparing alternative repayment contracts, I want to abstract from this motive, which could be accomplished with other tools (e.g., taxation). Second, given $\beta < R^{-1}$, this choice of discount rate is higher than the risk-free rate, consistent with the fact that student loan repayments likely have some correlation with aggregate shocks. In my baseline model, the average value of \mathcal{R}_a for $a \in (a_0, a_R)$ is 1.03. In Section 6.7, I consider the effect of using alternative discount rates.

²⁹I ignore the retirement pension because I remove means-testing based on wealth in all counterfactuals for reasons discussed in Appendix G.

³⁰I define the government budget in present-value terms rather than at the model's stationary distribution because the former has a more intuitive interpretation: it corresponds to the valuation implied by the first-order condition of a hypothetical lender with discount rate, \mathcal{R}_a . Additionally, this definition is preferable when considering budget-neutral repayment policies in subsequent analyses because it ensures a reasonable path for budget deficits in the transition between two policies without the difficulties associated with fully characterizing transition dynamics. In particular, this definition ensures that if the government were to immediately start giving loans to people graduating from college under two policies with equal values of \mathcal{G} , there would be no change in expected costs to this group of individuals.

The comparison of different repayment contracts in my model is contingent on the tax and transfer system, which is an alternative way to redistribute within- and across-individuals. For my normative analysis, I adopt the parametric specification of the tax system studied in [Heathcote, Storesletten, and Violante \(2017\)](#) calibrated to match the ATO Tax Schedule used during estimation. This specification is smooth and provides a close approximation to unconstrained optimal policies ([Heathcote and Tsujiyama 2021](#)), which is unlikely to be the case for the actual ATO Schedule. I also adopt a smoothed specification of the ATO unemployment benefit formula used in estimation; see Appendix [G](#) for additional details.

Results. The left panel of [Figure 13](#) presents the welfare and fiscal impact of the income-contingent loans used in the US and Australia. For clarity, this panel breaks the fiscal impact into the present value of the change in repayments and the change in other components \mathcal{G} , which are taxes net of transfers. To measure the dollar welfare impact of an alternative repayment contract, I compute the equivalent variation at $a = a_0$, which answers the following question:³¹ *“What value of a cash transfer at age a_0 would make an individual that attends college, prior to knowing her other initial states, indifferent between repaying under a new policy and repaying under 25-year fixed repayment?”*

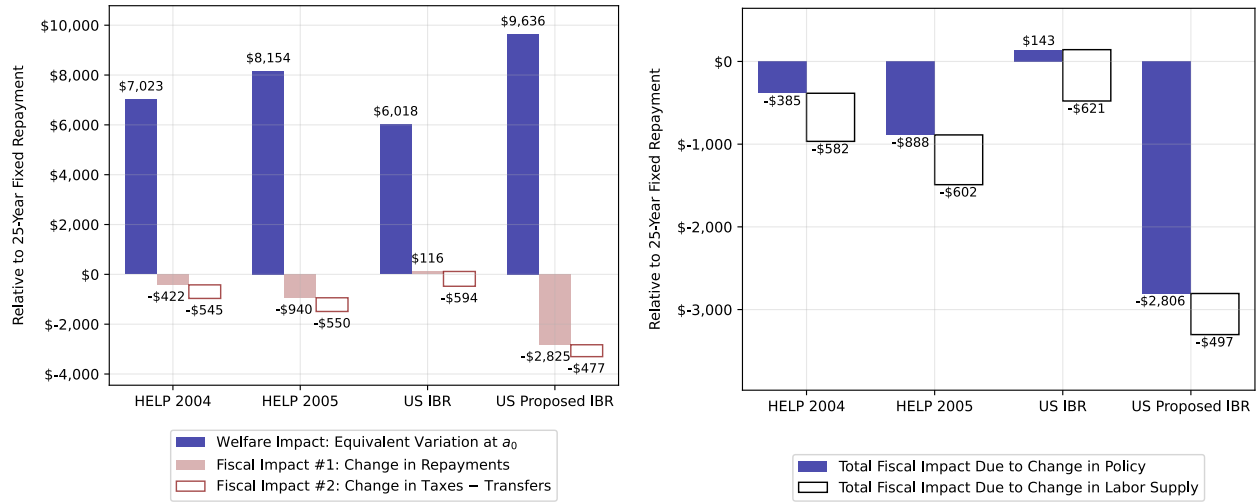
The first two columns show that both HELP repayment policies – before and after the policy change – provide gains equivalent to cash transfers of around \$7,500, which is 43% of the average initial debt balance in the model of \$17,500. These gains, however, come at a fiscal cost: in present value terms, the government collects around \$750 less in student loan repayments and \$550 in taxes net of transfers. The following two columns show the results for the income-based repayment (IBR) contract currently used in the US and the new IBR contract proposed by the Biden administration (known as “SAVE”), in which individuals make repay a fixed rate of income earned above a certain threshold.³² These two columns show that the two US IBR contracts deliver similar gains to HELP contracts but differ in fiscal cost. The current US IBR program has a fiscal cost that is around 60% lower than the HELP contracts, which is driven by the fact that repayments start at a lower value of income. In contrast, the proposed IBR program has a fiscal cost that is three times as large, reflecting the higher repayment threshold and lower repayment rate in this policy. Dividing the welfare gains by total fiscal cost delivers a marginal value of public funds (MVPF) for each policy relative to 25-year fixed repayment ([Finkelstein and Hendren 2020](#)). This MVPF (reported in [Figure A25](#)) is highest for US IBR and is high relative to typical estimates for policies targeting adults ([Hendren and Sprung-Keyser 2020](#)).

The right panel of [Figure 13](#) decomposes the total fiscal cost associated with moving from 25-year fixed repayment to income-contingent loans, the sum of the two fiscal impacts shown in the left panel, into two components. The first component, shown in the top black component of each bar,

³¹See Appendix [J](#) for details on the calculation of this welfare metric.

³²I implement these contracts without loan forgiveness in order for them to be comparable to the HELP contracts and return to the effects of forgiveness later.

Figure 13. Effects of Moving from Fixed Repayment to Existing Income-Contingent Loans



is the change in \mathcal{G} holding fixed individuals' labor supply decisions at their values under 25-year fixed repayment. The second component, shown below in blue, is the incremental change in \mathcal{G} due to the endogenous adjustment of labor supply. In other words, this second component measures the additional cost of the moral hazard created by income-contingent loans and would be zero in a model with exogenous labor supply. This moral hazard accounts for around 50% of the total cost from switching to HELP 2004 and HELP 2005 and 130% for US IBR. For US Proposed IBR, it accounts for only 15% of the fiscal cost, reflecting that the smaller 5% repayment rate generates a smaller behavioral response than the 10% rate under US IBR.

Effects of changing labor supply elasticity. Figure A26 reproduces the right panel of Figure 13 for different values of the elasticity of labor supply, ϕ . Increasing ϕ to twice its estimated value leads to a doubling of the cost of moral hazard, while reducing it by half leads to a cost reduction of over 60%. These results highlight the importance of correctly identifying the labor supply elasticity for quantifying the fiscal impact of income-contingent loans.

6.2 A Constrained-Optimal Income-Contingent Loan Provides Welfare Gains at Same Fiscal Cost as Fixed Repayment

The evidence in Section 6.1 shows that the welfare gains of existing income-contingent loans are large relative to their fiscal costs. In this section, I construct income-contingent loans with the same fiscal costs as fixed repayment contracts to assess whether these contracts can still provide gains and study the optimal form of income-contingent loans at a given cost.

Definition of constrained-planner's problem. I consider a social planner that maximizes borrower welfare by choosing one mandatory repayment contract. I assume this planner is constrained

à la [Ramsey \(1927\)](#) to choosing income-contingent loans with the same structure as US IBR contracts and the income-contingent loans used in the UK. These contracts have two parameters that make them essentially call options on individuals' incomes: the threshold at which repayment begins, K , and a repayment rate of income above the threshold, ψ . Aside from tractability, this restriction of the contract space is motivated by practical constraints that make implementing [Mirlees \(1974\)](#)-style optimal policies difficult ([Piketty and Saez 2013](#)).

The social planner's problem is thus:

$$\max_{\{\psi, K\}} \mathbf{E}_0 \left(V_{ia_0}^{1-\gamma} \right)^{\frac{1}{1-\gamma}}, \quad (17)$$

subject to:

$$\begin{aligned} \mathbf{E}_0 \left(\sum_{a=a_0}^{a_T} \frac{\tau_{ia} - u_{ia} - c_{ia}}{\mathcal{R}_a} + \frac{d_{ia}}{\mathcal{R}_a} - D_{ia_0} \right) &\geq \bar{\mathcal{G}}, \\ d_{ia} &= \min \{ \psi * \max \{ y_{ia} - K, 0 \}, D_{ia} \} * \mathbf{1}_{a \leq a_R}, \\ \psi &\in [0, 1], \quad K \geq 0. \end{aligned}$$

The objective function in this problem corresponds to the Epstein-Zin certainty equivalent functional of the stochastic consumption and labor supply streams, which depends (implicitly) on the three policy parameters, to an individual who is “behind the veil of ignorance” with respect to her initial conditions.³³ The first constraint requires that the fiscal revenue from the chosen repayment contract is at least $\bar{\mathcal{G}}$. I set $\bar{\mathcal{G}}$ equal to the revenue raised from a 25-year fixed repayment contract, which serves as a natural benchmark to an income-contingent loan given it has a similar repayment duration and is currently available to borrowers in the US. The second and third constraints capture the informational and parametric restrictions on the repayment contract.

Solving (17) is numerically challenging, especially when considering higher-dimensional contracts, because it is a nonlinear constrained optimization problem in which the objective and constraints do not have closed-forms. I thus leverage a combination of barrier methods in numerical optimization ([Nocedal and Wright 2006](#)) and the TikTak global optimizer from [Arnoud et al. \(2019\)](#) detailed in Appendix K.


Solution to planner's problem. The red solid line in [Figure 14](#) plots repayments as a function of income on the constrained-optimal income-contingent loan that solves (17) for an individual with a median initial debt balance. This contract provides individuals with significant insurance relative to a fixed repayment contract, as payments do not start until the 26th percentile of the income distribution at $K = \$27,147$. This value of K is similar to the threshold at which repayments begin in HELP 2004 system but lower than in HELP 2005. In US IBR contracts, K is set equal to 1.5 times the US federal poverty line, which corresponds to $1.5 * \$12,320 = \$18,480$ in 2005 AUD, or 68% of

³³As in Section 6.1, I focus only on the welfare of college-educated individuals with $\mathcal{E}_i = 1$.

the optimal value of K .³⁴

In order to collect sufficient revenue with a relatively high repayment threshold, the constrained-optimal contract has a repayment rate of $\psi = 33\%$, around three times the 10% repayment rate on current US IBR contracts. In other words, the optimal contract provides more insurance than current US IBR contracts by reducing payments from low-income borrowers in exchange for payments from high-income borrowers, with repayments are capped by initial debt balances. Although this repayment rate is relatively high, it induces almost no bunching at the repayment threshold, as shown in the income distribution in blue. This lack of bunching relative to the evidence in Figure 3 reflects the fact that this threshold changes the *marginal* rather than *average* repayment rate. As a result, individuals below the threshold do not receive an increase in their cash on hand, eliminating the liquidity effect discussed in Section 5.1. The small amount of bunching is consistent with the findings in Britton and Gruber (2020), who find limited bunching in the UK, where the repayment threshold changes the marginal repayment rate.

Figure 14. Structure of Constrained-Optimal Income-Contingent Loans



Effect of moral hazard. To isolate the impact of moral hazard on the design of income-contingent loans, the dashed green line in Figure 14 plots the contract that solves (17) in a model where individuals' labor supply is fixed at its value under the baseline 25-year fixed repayment contract.³⁵ The results show this alternative contract provides even more insurance than the contract in my baseline model, with a 30pp higher repayment rate and 40% higher threshold. This reflects


³⁴As of 2023, the US federal poverty line for a single household is \$14,580 USD. Deflating this to 2005 USD using the CPI and then converting to 2005 AUD using the USD/AUD exchange rate as of June 2005 delivers \$12,320. This value of the poverty line is similar to the value reported by the Melbourne Institute in 2005 of \$11,511.

³⁵In this model, I exclude disutility from labor supply from welfare since individuals are not choosing it.

that labor supply responses create a fiscal externality from a wedge between social and private incentives: individuals do not internalize that locating below the threshold reduces government revenue and affects the repayment contract the planner offers in equilibrium. Since the planner cannot raise a sufficient amount of revenue implementing the alternative contract in the baseline model because individuals reduce their labor supply, the planner lowers the repayment threshold to collect revenue from more individuals and the repayment rate to induce a smaller behavioral response.

Welfare gains. Figure 15 shows the welfare gain from the constrained-optimal income-contingent loan in Figure 14. To measure welfare gains, I use a consumption-equivalent metric commonly used by existing literature (e.g., Benabou 2002; Lucas 2003). This metric answers the following question: “What value of g would make an individual that attends college, prior to knowing her initial states, indifferent between repaying under a new policy and repaying under 25-year fixed repayment contract with their consumption increased by $g\%$ in every state of their life?” The leftmost blue bar in Figure 15 shows that the optimal income-contingent loan provides gains equivalent to a 1.32% increase in lifetime consumption relative to 25-year fixed repayment. This corresponds to 47% of the gain from forgiving debt balances entirely (which is not budget-neutral).

Figure 15. Decomposition of Welfare Gain from Constrained-Optimal Income-Contingent Loan



The right two bars in Figure 15 decompose the total gain shown in the first bar into the gain from providing insurance and loss from moral hazard. To compute the former, I solve (17) again instead assuming debt repayments, d_{ia} , can be made contingent on wage rates, w_{ia} , instead of income, y_{ia} . This contract is informationally-infeasible, but its gains depend entirely on the insurance benefits and not on labor supply responses. Therefore, the welfare cost of moral hazard corresponds to the difference between the gain of this wage-contingent loan and the constrained-optimal income-contingent loan. The results show that the cost of moral hazard is relatively small, accounting for

0.17pp or a 13% reduction in the total gain.

As discussed above, the repayment rate on the constrained-optimal income-contingent loan is higher than those in the US. [Figure A27](#) shows that imposing the constraint that $\psi \leq 10\%$, the current repayment rate on US IBR contracts, reduces the total gain by 0.20pp or 14%. Around half of this loss comes from a lower repayment rate requiring a lower repayment threshold to satisfy the government budget constraint, reducing the amount of insurance. The remaining half comes from the lower repayment threshold inducing labor supply responses by more individuals, which increases the loss from moral hazard.

6.3 Adding Forgiveness Reduces Welfare Gains of Income-Contingent Loans

I next consider the effects of adding forgiveness after a fixed horizon to income-contingent loans, as in the US and UK. The left panel of [Figure 16](#) compares the gain from the income-contingent loan in [Figure 14](#) with a constrained-optimal income-contingent loan with forgiveness at $a_0 + 20$, as in the currently available US IBR contracts. The latter contract generates a welfare gain of 0.99%, around 0.33pp lower than the contract without forgiveness repeated in the first column.³⁶


The reduction in the gain from adding forgiveness reflects a combination of two forces. First, adding forgiveness at the same fiscal cost requires a lower repayment threshold of $K = \$21,131$, as shown in the right panel of [Figure 16](#). The consequence of a lower repayment threshold is greater repayments from young borrowers in exchange for lowering repayments on older borrowers, for whom repayments are forgiven ([Figure A28](#)). This reduces the gain from insurance because younger borrowers have a higher marginal value of wealth from tighter borrowing constraints and stronger precautionary saving motives ([Gourinchas and Parker 2002](#); [Boutros et al. 2022](#)). The second force that reduces the gain from forgiveness is that a finite forgiveness horizon increases the loss from moral hazard ([Figure A29](#)). With a finite forgiveness horizon, borrowers are more willing to reduce their labor supply to lower repayments because it is less likely they will have to make these repayments later in their life cycle.

6.4 Income-Contingent Repayment Provides Larger Gains than Adding Forbearance

Relative to a fixed repayment contract, the income-contingent loan in [Figure 14](#) has two key differences. First, the latter contract provides payment reductions for low-income borrowers, whose income is below the repayment threshold. Second, income-contingent loans collect more payments from high-income borrowers, while repayments are independent of income under a fixed repayment contract. In reality, the fixed repayment contracts implemented in the US allow payments to

³⁶In untabulated results, I solve (17) optimizing over the forgiveness horizon and find no forgiveness is optimal.

Figure 16. Effects of Adding Forgiveness to Constrained-Optimal Income-Contingent Loans



be delayed if borrowers receive deferment, forbearance, or default. For example, [Figure A30](#) shows that 30% of outstanding debt in the US, as of 2019, was in one of these three non-repayment states.

[Figure 17](#) shows the welfare gain from a fixed repayment contract with forbearance, in which borrowers make constant payments that are independent of their income when their income is above \$16,384, the point at which unemployment benefits stop being paid in Australia (43% above the poverty line), but are allowed to access unlimited forbearance and make zero payments when their income falls below this point. The constant payment made outside of forbearance is calculated by solving (17) to ensure this contract has the same fiscal cost as other repayment contracts.³⁷ The left panel of [Figure 17](#) shows this contract delivers a gain of only 0.55% relative to 25-year fixed repayment, less than half of the gain from the constrained-optimal income-contingent loan.


The smaller gains from a fixed repayment contract with forbearance reflects the benefits of the call option-like structure of fully income-contingent repayment. As shown in the right panel of [Figure 17](#), the income-contingent loan collects more payments from high-income individuals. Although these individuals are likely to pay off their debt, the acceleration of these payments increases their expected discounted value. This in turn enables the social planner to have a higher repayment threshold, increasing insurance at a given cost.

6.5 Equity Contracts Generate Larger, but More Dispersed, Welfare Gains

The development of income-contingent loans was motivated by [Friedman \(1955\)](#), who advocated using equity contracts known as income-sharing agreements (ISAs) in which individuals re-

³⁷An alternative contract would be a 25-year fixed repayment contract with the same unemployment forbearance, where the interest rate is adjusted to balance the government budget. I choose not to use this contract to ensure close comparability with income-contingent loans: the former makes payments conditional on debt balances, while income-contingent loans do not. Nevertheless, this alternative contract delivers a similar gain of 0.52%.

Figure 17. Welfare Gains from Alternative Contracts: Forbearance and Equity Contracts



pay a percentage of their income for a fixed repayment period. Successful implementation of these contracts has been relatively limited due to adverse selection (Herbst and Hendren 2021). Setting aside these selection issues, I next use my model to assess the desirability of income-sharing agreements as a mandatory government-provided financing contract. I model ISAs after those provided by Purdue University in 2016-2017, in which individuals repay a constant fraction of their income for nine years (Mumford 2022).


The third column in the left panel of Figure 17 shows the welfare gain from a nine-year ISA is 0.82%, where the parameter controlling the share of income repaid has been adjusted to balance the government budget. This gain is around 40% (or 0.5pp) lower than the gain from the constrained-optimal income-contingent loan. The lower gain reflects the same force that generates smaller gains from forgiveness: a pure ISA requires payments from all borrowers in the first few years of their life, which is when they value repayment reductions the most, in exchange for zero payments when they are older (i.e., after nine years).

The final column of Figure 17 shows that a modified ISA offered by Purdue University does significantly better. In the 9-Year ISA with Threshold, borrowers only make payments when their income exceeds a certain threshold, which is chosen jointly with the income-share rate to solve (17). This contract performs better than a pure ISA because it avoids requiring payments from low-income young borrowers. Additionally, it outperforms the constrained-optimal income-contingent loan because it provides greater insurance. With an income-contingent loan, repayments from high-income individuals are capped by their initial debt balances. However, with an income-sharing agreement, these payments are uncapped and thus can be used to finance even lower repayments from low-income individuals. The right panel of Figure 17 shows how this manifests in a 70% higher repayment threshold of $K = \$46,821$.

Although equity contracts generate larger gains on average, these gains are also more heterogeneous. Figure 18 plots the distribution of welfare gains and losses at $a = a_0$ from the constrained-

optimal income-contingent loan and 9-Year ISA with Threshold. Relative to 25-year fixed repayment, the income-contingent loan gives gains for around 70% of individuals, while the remaining 30% experience small losses. These small losses are concentrated among high-income individuals, who are required to repay their loans faster under income-contingent relative to fixed repayment.

Figure 18. Distribution of Gains from Constrained-Optimal Income-Contingent Loan and Equity Contract



Notes: Due to computational constraints, the welfare gains plotted in this plot are computed as the percent change in initial value functions. This has the interpretation of a *net* consumption-equivalent welfare gain: it corresponds to the value of g that would make an individual indifferent between the original contract and having their consumption net of the utility of labor supply increased by $g\%$ in every state of their life. This quantity is quantitatively very similar to the consumption-equivalent welfare gain used in other plots, but is significantly easier to compute because it does not require solving a numerical fixed point for each initial state.

Figure 18 shows that the gains from the optimal equity contract are significantly more dispersed. Relative to 25-year fixed repayment, 18% of borrowers have losses greater than 0.5%, while only 1.2% of borrowers have losses this large under the income-contingent loan. This heterogeneity is primarily driven by losses from high-income individuals whose repayments are uncapped under an equity contract. Additionally, income-sharing agreements induce significant redistribution from individuals with low to individuals with high debt balances. The redistribution across debt balances is sufficiently large that it reverses the ranking of certainty-equivalent values across terciles of initial debt (Figure A31).

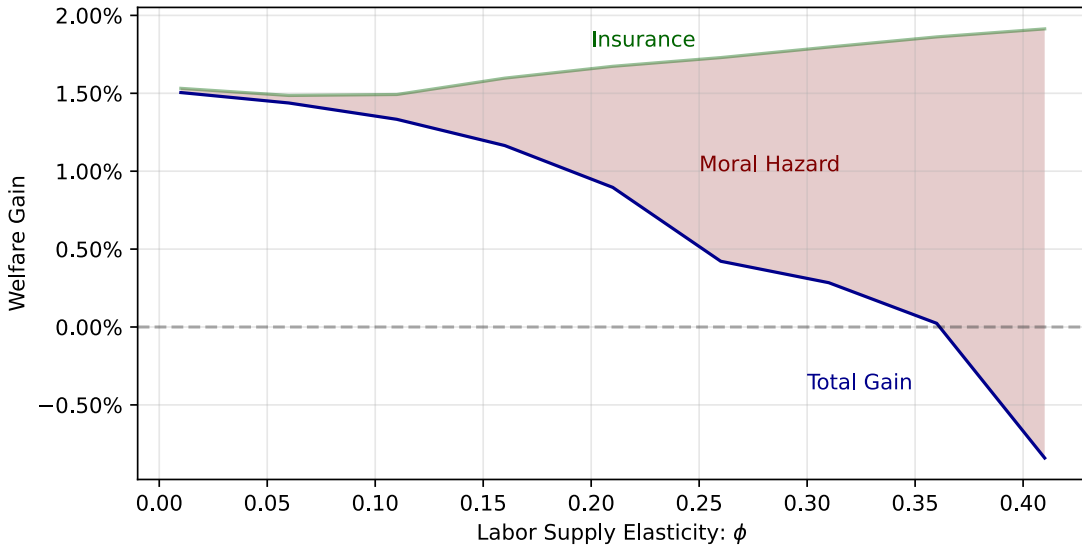
In sum, although (properly-designed) equity contracts give larger gains, they are also more likely to generate ex-ante responses not captured by my model because the distribution of gains across initial debt balances is significantly more dispersed. This suggests that income-contingent loans may be a more robust mechanism for implementing income-contingent repayment.

6.6 A Higher Labor Supply Elasticity Increases Welfare Loss from Moral Hazard, But These Losses Can Be Reduced with Alternative Contracts

To assess the robustness of the welfare gains from income-contingent repayment, I re-compute the income-contingent loan that solves (17) for different values of the labor supply elasticity, ϕ . The gains from the resulting constrained-optimal contracts are plotted in Figure 19. The results show that for income-contingent loans to deliver a welfare *loss* relative to fixed repayment contracts, ϕ would have to be above 0.37. This higher value of ϕ increases the moral hazard cost of income-contingent repayment while having minimal effect on its insurance benefits. At this value of ϕ , the loss from moral hazard is 10 times as large as the baseline model and outweighs the insurance benefits.

Figure A32 plots the fit of this alternative model in which fixed repayment is optimal on the most important set of moments for identifying the labor supply elasticity in structural estimation: bunching around the repayment thresholds. These results show that this model generates significantly more bunching than the baseline model and the data. Quantitatively, the number of individuals below relative to above the threshold is around 70% larger, more than within any occupation.

Figure 19. Welfare Gains from Constrained-Optimal Income-Contingent Loan as a Function of ϕ




Motivated by Figure 19, I next consider whether the following alternative forms of income-contingent loans can reduce the welfare cost of moral hazard, even when $\phi = 0.37$:

$$\begin{aligned} \text{Smooth Income-Contingent Loan : } d_{ia} &= \min \left\{ \max \left\{ \psi_0 + \psi_1 y_{ia} + \psi_2 y_{ia}^2, 0 \right\}, D_{ia} \right\}, \\ \text{Income-Contingent Loan + Age : } d_{ia} &= \min \left\{ \max \left\{ \psi_0 + \psi_1 y_{ia} + \psi_2 y_{ia}^2 + \psi_3 a, 0 \right\}, D_{ia} \right\}, \\ \text{Income-Contingent Loan + Debt : } d_{ia} &= \min \left\{ \max \left\{ \psi_0 + \psi_1 y_{ia} + \psi_2 y_{ia}^2 + \psi_3 D_{ia}, 0 \right\}, D_{ia} \right\}. \end{aligned}$$

The first contract corresponds to a smoothed version of the US IBR-style income-contingent loans considered above, in which repayments are a quadratic function of income. The latter two contracts make repayments conditional on age and debt, respectively. For each of these alternative contracts, I solve (17) to find the constrained-optimal values of $\{\psi_i\}$.

Figure 20 shows the gains from these contracts in my baseline model with $\phi = 0.37$, in which the US IBR-style income-contingent loan delivers losses relative to 25-year fixed repayment.³⁸ The results show that all three of these contracts restore the gains of income-contingent repayment. These large gains come entirely from reducing the cost of moral hazard: the smooth income-contingent loan generates a 105pp reduction in this cost, while the age- and debt-contingent contracts generate additional 32pp and 25pp reductions, respectively. Figure A34 plots these constrained-optimal repayment contracts, which shows the smooth income-contingent loan takes a similar shape to the IBR-style loan. However, the smoother repayment structure helps minimize labor supply responses, reducing the cost of moral hazard. The age- and debt-contingent contracts increase payments with age and debt balances, both of which further reduce the moral hazard from income-contingent repayment by increasing the future cost associated with reducing current labor supply.

Figure 20. Welfare Gains from Alternative Forms of Income-Contingent Loans when $\phi = 0.37$



6.7 Sensitivity of Welfare Gains to Model Misspecification

This section assesses the robustness of the welfare gains from the constrained-optimal income-contingent loans that solve (17) to model misspecification by considering several extensions and alternative parameter values. The results are presented in the rows in Table 6 and discussed below.

³⁸Figure A33 shows the results in the baseline model. In this model, all three contracts deliver almost identical gains to the income-contingent loan considered in previous sections. This result reflects that the discontinuity in the marginal repayment rate at the repayment threshold is not very distortionary in the baseline model, as evident from the lack of bunching in Figure 14.

Risk and time preferences. Rows (1)-(3) of [Table 6](#) shows the effect of moving the coefficient of relative risk aversion, γ , and the elasticity of intertemporal substitution, σ^{-1} . Starting from the baseline values, I first set $\gamma = 7.5$ as in [Bansal and Yaron \(2004\)](#), which increases both individuals' risk aversion but also introduces a demand for the early resolution of uncertainty ([Epstein and Zin 1989](#)). This increases the gain from income-contingent repayment, as individuals value the reductions in repayments more. The optimal contract also provides significantly more insurance. Next, I set $\sigma^{-1} = 1.5$ as in [Bansal and Yaron \(2004\)](#). This reduces the gain but has a relatively small effect on the optimal contract.

Level, uncertainty, and redistributive effects. In general, the consumption-equivalent welfare gain of a policy reform is the sum of three effects: (i) level effects due to changes in average consumption, (ii) uncertainty effects due to changes in the volatility of the agents' consumption paths that affects welfare because of risk aversion and incomplete markets, and (iii) redistributive effects due to changes in consumption-equivalents across different initial conditions ([Benabou 2002](#)). Due to the non-homotheticity and non-convexities in my model, calculating these terms analytically is not possible.³⁹ To assess the importance of each of these terms, I compare the gain in the baseline model with the gains in two alternative models: a model without any ex-post uncertainty (aside from Calvo shocks) and a model without any ex-ante and ex-post uncertainty. Intuitively, the gain of the latter model should be due to level effects, while the difference between the two captures redistributive effects. The size of uncertainty effects can be estimated by comparing the baseline model to the model with no ex-post uncertainty. The results from these two models are shown in rows (4) and (5), which show that most of the gain comes from redistributive and uncertainty effects, with around half coming from each.

Occupation-level heterogeneity. [Figure 4](#) shows a significant amount of occupation-level heterogeneity in labor supply responses that is not in my baseline model. To assess the importance of this heterogeneity, I consider an extension in my model where there are two types of educated individuals with different values of the Calvo parameter, λ . To assess robustness with respect to an extreme amount of heterogeneity, I assume these groups have $\lambda = 0$ and $\lambda = 1$ respectively. I then calibrate the fraction of each type so that the model generates the same fit on the amount of bunching around the 2005 0% repayment threshold. Row (6) shows that incorporating this heterogeneity decreases the gain, but the overall gain is still positive.

Learning-by-doing. Row (7) shows the results from the model with learning-by-doing estimated in column (5) of [Table 4](#). This model generates similar but slightly larger gains than the baseline model. This reflects that, with GHH preferences, individuals have lower labor supply early in life

³⁹An alternative decomposition that can be implemented numerically is presented in [Abbott et al. \(2019\)](#). I cannot perform this decomposition in my model because I do not discretize initial states. This would require me to calculate consumption-equivalents for many states, which is not computationally feasible since each consumption-equivalent requires solving a fixed point. My approach avoids this problem but lacks analytical results guaranteeing it generates accurate estimates of the three effects.

Table 6. Robustness of Welfare Gains from Constrained-Optimal Income-Contingent Loans

	Deviation from Baseline	Welfare Gain	= Insurance	+ Moral Hazard	ψ^*	K^*
(1)	RRA = 7.5	3.52%	4.00%	-0.48%	50%	\$27,607
(2)	EIS = 1.5	0.57%	0.7%	-0.13%	42%	\$30,905
(3)	RRA = 7.5, EIS = 1.5	1.87%	2.29%	-0.43%	49%	\$28,641
(4)	No Ex-Post Uncertainty	0.58%	0.76%	-0.17%	27%	\$18,098
(5)	No Uncertainty	-0.17%	0.15%	-0.32%	21%	\$26,906
(6)	Occupation Heterogeneity	0.28%	0.33%	-0.05%	22%	\$25,639
(7)	Learning-by-Doing	1.68%	.	.	35%	\$36,615
(8)	Wealth Effects on Labor Supply	0.82%	1.05%	-0.23%	37%	\$30,307
(9)	$\rho = 0.8$	0.90%	1.14%	-0.23%	42%	\$34,244
(10)	$\rho = 0.99$	1.35%	1.63%	-0.28%	35%	\$18,949
(11)	Non-Normal Permanent Shocks	1.14%	1.43%	-0.30%	28%	\$26,933
(12)	$r_d = 2\%$	1.96%	2.14%	-0.18%	38%	\$47,731
(13)	US Tax System	1.18%	1.36%	-0.19%	38%	\$28,838
(14)	$\mathcal{R}_a = R$	1.06%	1.41%	-0.35%	29%	\$22,696
(15)	$\mathcal{R}_a = R + 4\%$	1.60%	1.65%	-0.05%	46%	\$34,441
	Baseline Model	1.32%	1.47%	-0.15%	33%	\$27,147

Notes: The welfare gain from insurance and welfare cost of moral hazard are not reported for the learning-by-doing model because in that model wage rates are endogenous, so a wage-contingent repayment contract also introduces some welfare costs of moral hazard.

under a fixed repayment contract when their consumption is low. With learning-by-doing, there is an added benefit of increasing labor supply early in life because it delivers higher income and thus greater revenue later in life. This effect is larger than the second effect that learning-by-doing introduces, in which labor supply reductions to avoid repayments produce longer-run costs.

Wealth effects on labor supply. There is disagreement in existing literature on the size of wealth effects on labor supply. For example, [Cesarini et al. \(2017\)](#) find relatively small wealth effects from lottery winnings in Sweden, while [Golosov, Graber, Mogstad, and Novgorodsky \(2023\)](#) find larger effects from lottery winnings in the US. To assess the importance of wealth effects, I adjust the flow utility in (5) to be:

$$\frac{1}{\eta} \left(\frac{c_{ia}}{n_a} \right)^\eta - \kappa \frac{\ell_{ia}^{1+\phi^{-1}}}{1+\phi^{-1}}.$$

I follow [Keane \(2011\)](#) and consider $\eta = 0.5$. Row (8) shows this has a minimal effect on the optimal contract and reduces the gain slightly.

Persistence of income risk. Rows (9) and (10) consider the effect of alternative values of ρ . As expected, higher values of ρ , which increase the quantity of risk against which individuals would like to insure, raise the gains from income-contingent repayment. Lower values of ρ do the opposite.

Non-normal income risk. A recent body of empirical evidence using administrative datasets on individuals' incomes highlights the importance of non-normal income shocks ([Guvenen et al. 2014, 2021](#); [Braxton et al. 2021](#)), which has been shown to have implications for portfolio choice ([Catherine 2022](#)), asset prices ([Schmidt 2016](#)), and mortgage contract design ([Campbell et al.](#)

2021). I introduce this into my baseline model, without re-estimating the model, by making ν_{ia} in (7) be drawn from a mixture of independent two normal distributions with different means and variances. I calibrate these additional parameters by estimating two models with exogenous income processes with and without the mixture of normals. I then multiply the values of the parameters in the former by the ratio of my estimates in Table 4 to the latter estimates. Row (11) shows this has a small effect on my baseline results.

Interest rate on debt. In my baseline analysis, I set the real interest rate on debt balances to zero, as in the HELP program. However, in the US, debt balances are subject to interest accumulation. As a result, I consider an alternative interest rate of 2% above the real interest rate, which is similar to the markup of student loans above Treasury bill rates in the US (Ji 2021) and above the Bank of England base rate in the UK (Britton and Gruber 2020). Row (12) shows this gives a larger gain, reflecting that there is more room to provide insurance when interest payments can be collected from borrowers who would have paid off their debts.

Alternative tax system. My analysis is contingent on the current tax and transfer system in the model because student debt policies may be trying to undo suboptimalities in tax system. To assess the robustness of the results, I recompute my gains using the tax and transfer system from Heathcote et al. (2017) that approximates the US system. The results in Row (13) show similar gains, but the optimal contract provides more insurance to account for the US tax system being less progressive.

Discount rates for the government budget. My model does not have aggregate risk, so the proper discount rate for debt repayments is the risk-free rate. However, in reality, student loan repayments likely should be discounted at a higher rate, given that they are income-dependent and thus are correlated with the business cycle. To address this, I consider two alternative discount rates: the risk-free rate and the risk-free rate plus a constant 4% risk premium. Rows (14) and (15) a lower (higher) discount rate decreases (increases) the gains.

7 Conclusion

This paper studies the trade-off between providing insurance and disincentivizing labor supply in student loans with income-contingent repayment. Using discontinuities in repayment rates from Australia’s student loan system, I show borrowers adjust their labor supply to reduce repayments on income-contingent loans. These responses are larger in occupations with more hourly flexibility, among young borrowers with more debt, and among liquidity-constrained borrowers with less wealth and larger housing payments. I use these responses to estimate a structural model and find they are consistent with a Frisch labor supply elasticity of 0.11 and substantial frictions that limit labor supply adjustment. In this model, a constrained-optimal income-contingent loan generates

welfare gains relative to a fixed repayment contract equivalent to a 1.3% increase in lifetime consumption, with the same fiscal cost. Equity contracts generate gains that are larger on average but significantly more dispersed. The labor supply responses to income-contingent repayment reduce the insurance these contracts can provide at a given cost, but they are too small to justify fixed repayment contracts.

References

- Abbott, Brant, Giovanni Gallipoli, Costas Meghir, and Giovanni L. Violante (2019), “Education policy and intergenerational transfers in equilibrium.” *Journal of Political Economy*, 127, 2569–2624.
- Abel, Andrew B., Janice Eberly, and Stavros Panageas (2013), “Optimal Inattention to the Stock Market With Information Costs and Transactions Costs.” *Econometrica*, 81, 1455–1481.
- Abraham, Katharine G., Emel Filiz-Ozbay, Erkut Y. Ozbay, and Lesley J. Turner (2020), “Framing effects, earnings expectations, and the design of student loan repayment schemes.” *Journal of Public Economics*, 183.
- Amromin, Gene and Janice Eberly (2016), “Education Financing and Student Lending.” *Annual Review of Financial Economics*, 8, 289–315.
- Andersen, Steffen, John Y. Campbell, Kasper Meisner Nielsen, and Tarun Ramadorai (2020), “Sources of inaction in household finance: Evidence from the danish mortgage market.” *American Economic Review*, 110, 3184–3230.
- Arnoud, Antoine, Fatih Guvenen, and Tatjana Kleineberg (2019), “Benchmarking Global Optimizers.” *Working Paper*, w26340.
- Auclert, Adrien, Will S Dobbie, and Paul S. Goldsmith-Pinkham (2019), “Macroeconomic Effects of Debt Relief: Consumer Bankruptcy Protections in the Great Recession.” *Working Paper*.
- Auclert, Adrien and Matthew Rognlie (2017), “A note on multipliers in NK models with GHH preferences.” *Working Paper*, 14.
- Australian Council of Social Service (2017), “Ending tax avoidance, evasion and money laundering through private trusts.” ACOSS Policy Briefing.
- Australian Council of Social Service (2018), “Components of Australia’s wealth.”
- Baily, Neil (1978), “Some Aspects of Optimal Unemployment Insurance.” *Journal of Public Economics*, 10, 379–402.
- Bansal, Ravi and Amir Yaron (2004), “Risks for the long run: A potential resolution of asset pricing puzzles.” *Journal of Finance*, 59, 1481–1509.
- Barr, Nicholas, Bruce Chapman, Lorraine Dearden, and Susan Dynarski (2019), “The US college loans system: Lessons from Australia and England.” *Economics of Education Review*, 71, 32–48.
- Beer, Gillian and Bruce Chapman (2004), “HECS System Changes: Impact on Students.” *Agenda - A Journal of Policy Analysis and Reform*, 11.
- Belley, Philippe and Lance Lochner (2007), “The Changing Role of Family Income and Ability in Determining Educational Achievement.” *Journal of Human Capital*, 1, 37–89.
- Benabou, Roland (2002), “Tax and Education Policy in a Heterogeneous-Agent Economy: What Levels of Redistribution Maximize Growth and Efficiency?” *Econometrica*, 70, 481–517.
- Best, Michael and Henrik Kleven (2012), “Optimal Income Taxation with Career Effects of Work Effort.” *Working Paper*.
- Black, Sandra E., Jeffrey T. Denning, Lisa J. Dettling, Sarena Goodman, and Lesley J. Turner (2022), “Taking It to the Limit: Effects of Increased Student Loan Availability on Attainment, Earnings, and Financial Well-Being.” *Working Paper*.
- Blundell, Richard, Luigi Pistaferri, and Ian Preston (2008), “Consumption Inequality and Partial Insurance.” *American Economic Review*, 98, 1887–1921.
- Bommier, Antoine, Daniel Harenberg, François Le Grand, and Cormac O’Dea (2020), “Recursive Preferences, the Value of Life, and Household Finance.” *SSRN Electronic Journal*.
- Boutros, Michael, Nuno Clara, and Francisco Gomes (2022), “Borrow Now, Pay Even Later: A Quantitative Analysis of Student Debt Payment Plans.” *SSRN Electronic Journal*.
- Bovenberg, A. Lans and Bas Jacobs (2005), “Redistribution and education subsidies are Siamese twins.” *Journal of Public Economics*, 89, 2005–2035.
- Braxton, J. Carter, Kyle Herkenhoff, Jonathan Rothbaum, and Lawrence D. W. Schmidt (2021), “Changing income risk across the US skill distribution: Evidence from a generalized Kalman filter.” *Working Paper*.

- Britton, Jack and Jonathan Gruber (2020), “Do income contingent student loans reduce labor supply?” *Economics of Education Review*, 79, 102061.
- Caballero, Ricardo J. and Eduardo M. R. A. Engel (1999), “Explaining Investment Dynamics in U.S. Manufacturing: A Generalized (S, s) Approach.” *Econometrica*, 67, 783–826.
- Calvo, Guillermo A. (1983), “Staggered prices in a utility-maximizing framework.” *Journal of Monetary Economics*, 12, 383–398.
- Campbell, John Y., Nuno Clara, and João F. Cocco (2021), “Structuring Mortgages for Macroeconomic Stability.” *The Journal of Finance*, 76, 2525–2576.
- Caplin, Andrew, James H. Carr, Fredrick Pollock, and Zhong Yi Tong (2007), “Shared-Equity Mortgages, Housing Affordability, and Homeownership.” *Working Paper*.
- Caplin, Andrew and Daniel F. Spulber (1987), “Menu Costs and the Neutrality of Money.” *Quarterly Journal of Economics*, 102, 703–726.
- Carneiro, Pedro and James J. Heckman (2002), “The Evidence on Credit Constraints in Post-Secondary Schooling.” *The Economic Journal*, 112, 705–734.
- Carroll, Christopher D (1997), “Buffer-Stock Saving and the Life-Cycle/Permanent Income Hypothesis.” *The Quarterly Journal of Economics*, 112, 12–26.
- Carroll, Christopher D. and Miles S. Kimball (1996), “On the Concavity of the Consumption Function.” *Econometrica*, 64, 981–992.
- Catherine, Sylvain (2022), “Countercyclical Income Risk and Portfolio Choices over the Life-Cycle.” *Review of Financial Studies*, 35, 4054.
- Catherine, Sylvain and Constantine Yannelis (2023), “The Distributional Effects of Student Loan Forgiveness.” *Journal of Financial Economics*.
- Caucutt, Elizabeth M. and Lance Lochner (2020), “Early and late human capital investments, borrowing constraints, and the family.” *Journal of Political Economy*, 128, 1065–1147.
- Cesarini, David, Erik Lindqvist, Matthew J. Notowidigdo, and Robert Ostling (2017), “The Effect of Wealth on Individual and Household Labor Supply: Evidence from Swedish Lotteries.” *American Economic Review*, 107, 34.
- Chakrabarti, Rajashri, Vyacheslav Fos, Andres Liberman, and Constantine Yannelis (2020), “Tuition, Debt, and Human Capital.” *Working Paper*.
- Chapman, Bruce (2006), “Chapter 25 Income Contingent Loans for Higher Education: International Reforms.” In *Handbook of the Economics of Education*, volume 2, 1435–1503, Elsevier.
- Chapman, Bruce and Andrew Leigh (2009), “Do Very High Tax Rates Induce Bunching? Implications for the Design of Income Contingent Loan Schemes.” *Economic Record*, 85, 276–289.
- Chapman, Bruce and Tony Salvage (2001), “Australian Postgraduate Financing Options.” *Agenda - A Journal of Policy Analysis and Reform*, 8.
- Chetty, R., J. N. Friedman, T. Olsen, and L. Pistaferri (2011), “Adjustment Costs, Firm Responses, and Micro vs. Macro Labor Supply Elasticities: Evidence from Danish Tax Records.” *The Quarterly Journal of Economics*, 126, 749–804.
- Chetty, Raj (2006), “A general formula for the optimal level of social insurance.” *Journal of Public Economics*, 90, 1879–1901.
- Chetty, Raj (2008), “Moral hazard versus liquidity and optimal unemployment insurance.” *Journal of Political Economy*, 116, 173–234.
- Chetty, Raj (2009), “Is the Taxable Income Elasticity Sufficient to Calculate Deadweight Loss? The Implications of Evasion and Avoidance.” *American Economic Journal: Economic Policy*, 1, 31–52.
- Chetty, Raj (2012), “Bounds on Elasticities With Optimization Frictions: A Synthesis of Micro and Macro Evidence on Labor Supply.” *Econometrica*, 80, 969–1018.
- Chetty, Raj and Amy Finkelstein (2013), “Social Insurance: Connecting Theory to Data.” In *Handbook of Public Economics*, volume 5, Elsevier B.V.
- Chetty, Raj, John N Friedman, and Emmanuel Saez (2013), “Using Differences in Knowledge Across Neighborhoods to Uncover the Impacts of the EITC on Earnings.” *American Economic Review*, 103, 2683–2721.

- Chetty, Raj, Adam Looney, and Kory Kroft (2009), “Salience and taxation: Theory and evidence.” *American Economic Review*, 99, 1145–1177.
- Choukhmane, Taha (2021), “Default Options and Retirement Saving Dynamics.” *Working Paper*, 1–78.
- Choukhmane, Taha and Tim de Silva (2023), “What Drives Investors’ Portfolio Choices? Separating Risk Preferences from Frictions.” *Working Paper*.
- Cocco, João F., Francisco J. Gomes, and Pascal J. Maenhout (2005), “Consumption and portfolio choice over the life cycle.” *Review of Financial Studies*, 18, 491–533.
- Cochrane, John H. (1991), “A simple test of consumption insurance.” *Journal of Political Economy*, 93, 957–978.
- Deaton, Angus and Christina H Paxson (1994), “Intertemporal Choice and Inequality.” *Journal of Political Economy*, 102, 437–468.
- DeFusco, Anthony A., Tang Huan, and Constantine Yannelis (2022), “Measuring the Welfare Cost of Asymmetric Information in Consumer Credit Markets.” *Working Paper*.
- Department of Education and Training (2023), “HECS-HELP Determinations.” Technical report, Australian Government.
- Di Maggio, Marco, Ankit Kalda, and Vincent W. Yao (2021), “Second Chance: Life without Student Debt.” *Journal of Finance*.
- Dobbie, Will and Jae Song (2015), “Debt Relief and Debtor Outcomes: Measuring the Effects of Consumer Bankruptcy Protection.” *American Economic Review*, 105, 1272–1311.
- Duffie, Darrell and Kenneth J. Singleton (1993), “Simulated Moments Estimation of Markov Models of Asset Prices.” *Econometrica*, 61, 929–952.
- D’Souza, Gabriela (2018), “A higher education bubble?” *Working Paper*.
- Ebrahimian, Mehran (2020), “Student Loans and Social Mobility.” *SSRN Electronic Journal*.
- Einav, Liran and Amy Finkelstein (2011), “Selection in Insurance Markets: Theory and Empirics in Pictures.” *Journal of Economic Perspectives*, 25, 115–138.
- Einav, Liran, Amy Finkelstein, and Paul Schrimpf (2017), “Bunching at the kink: Implications for spending responses to health insurance contracts.” *Journal of Public Economics*, 146, 27–40.
- Einav, Liran, Mark Jenkins, and Jonathan Levin (2012), “Contract Pricing in Consumer Credit Markets.” *Econometrica*, 80, 1387–1432.
- Epstein, Larry G. and Stanley E. Zin (1989), “Substitution, Risk Aversion, and the Temporal Behavior of Consumption and Asset Returns: A Theoretical Framework.” *Econometrica*, 57, 937.
- Ey, Carol (2021), “The Higher Education Loan Program (HELP) and related loans: A chronology.” Technical report, Parliament of Australia.
- Feldstein, Martin (1999), “Tax Avoidance and the Deadweight Loss of the Income Tax.” *Review of Economics and Statistics*, 81, 674–680.
- Finkelstein, Amy and Nathaniel Hendren (2020), “Welfare analysis meets causal inference.” *Journal of Economic Perspectives*, 34, 146–167.
- Folch, Marc and Luca Mazzone (2021), “Go Big or Buy a Home: Student Debt, Career Choices and Wealth Accumulation.” *SSRN Electronic Journal*.
- Friedman, Milton (1955), “The Role of Government in Education.” In *Economics and the Public Interest* (Robert A. Solow, ed.), Rutgers University Press, New Brunswick, New Jersey.
- Ganong, Peter and Pascal Noel (2019), “Consumer spending during unemployment: Positive and normative implications.” *American Economic Review*, 109, 2383–2424.
- Ganong, Peter and Pascal Noel (2020), “Liquidity versus wealth in household debt obligations: Evidence from housing policy in the great recession.” *American Economic Review*, 110, 3100–3138.
- Ganong, Peter and Pascal Noel (2022), “Why Do Borrowers Default on Mortgages?” *The Quarterly Journal of Economics*.
- Gervais, Martin, Qian Liu, and Lance Lochner (2022), “The Insurance Implications of Government Student Loan Repayment Schemes.” *Working Paper*.

- Giglio, Stefano, Matteo Maggiori, Johannes Stroebel, and Stephen P. Utkus (2021), “Five Facts About Beliefs and Portfolios.” *American Economic Review*, 115, 1481–1522.
- Golosov, Mikhail, Michael Graber, Magne Mogstad, and David Novgorodsky (2023), “How Americans Respond to Idiosyncratic and Exogenous Changes in Household Wealth and Unearned Income.” *Working Paper*.
- Gourinchas, Pierre-Olivier and Jonathan A. Parker (2002), “Consumption over the life cycle.” *Econometrica*, 70, 47–89.
- Greenwald, Daniel L., Tim Landvoigt, and Stijn Van Nieuwerburgh (2021), “Financial Fragility with SAM?” *Journal of Finance*, 76, 651–706.
- Greenwood, Jeremy, Zvi Hercowitz, and Gregory W Huffman (1988), “Investment, Capacity Utilization, and the Real Business Cycle.” *American Economic Review*, 78, 402–417.
- Gruber, Jonathan (1997), “The Consumption Smoothing Benefits of Unemployment Insurance.” *American Economic Review*, 87, 192–205.
- Gupta, Arpit and Christopher Hansman (2022), “Selection, Leverage, and Default in the Mortgage Market.” *The Review of Financial Studies*, 35, 720–770.
- Guvenen, Fatih (2009a), “An empirical investigation of labor income processes.” *Review of Economic Dynamics*, 12, 58–79.
- Guvenen, Fatih (2009b), “A Parsimonious Macroeconomic Model for Asset Pricing.” *Econometrica*, 77, 1711–1750.
- Guvenen, Fatih, Fatih Karahan, Serdar Ozkan, and Jae Song (2021), “What Do Data on Millions of U.S. Workers Reveal About Lifecycle Earnings Dynamics?” *Econometrica*, 89, 2303–2339.
- Guvenen, Fatih, Alisdair McKay, and Conor Ryan (2022), “A Tractable Income Process for Business Cycle Analysis.” *Working Paper*.
- Guvenen, Fatih, Serdar Ozkan, and Jae Song (2014), “The nature of countercyclical income risk.” *Journal of Political Economy*, 122, 621–660.
- Hampole, Menaka V (2022), “Financial Frictions and Human Capital Investments.” *Working Paper*.
- Hanson, Melanie (2022), “Student Loan Default Rate.” Technical report, Education Data Initiative.
- Hanson, Melanie (2023), “Average Cost of College & Tuition.” Technical report, Education Data Initiative.
- Hartman-Glaser, Barney and Benjamin Hébert (2020), “The Insurance Is the Lemon: Failing to Index Contracts.” *Journal of Finance*, 75, 463–506.
- Heathcote, Jonathan, Kjetil Storesletten, and Giovanni L. Violante (2017), “Optimal Tax Progressivity: An Analytical Framework.” *The Quarterly Journal of Economics*, 132, 1693–1754.
- Heathcote, Jonathan and Hitoshi Tsujiyama (2021), “Optimal Income Taxation: Mirrlees Meets Ramsey.” *Journal of Political Economy*, 129, 3141–3184.
- Hendren, Nathaniel and Ben Sprung-Keyser (2020), “A Unified Welfare Analysis of Government Policies.” *The Quarterly Journal of Economics*, 135, 1209–1318.
- Herbst, Daniel (2023), “Liquidity and Insurance in Student-Loan Contracts: The Effects of Income-Driven Repayment on Borrower Outcomes.” *American Economic Journal: Applied Economics*, 15, 1–25.
- Herbst, Daniel and Nathaniel Hendren (2021), “Opportunity Unraveled: Private Information and the Missing Markets for Financing Human Capital.” *Working Paper*.
- Herbst, Daniel, Miguel Palacios, and Constantine Yannelis (2023), “Equity and Incentives in Household Finance.” *Working Paper*.
- Huang, Yueling (2022), “Rethinking College Financing: Wealth, College Majors, and Macroeconomic Consequences.” *Working Paper*.
- Imai, Susumu and Michael P. Keane (2004), “Intertemporal Labor Supply and Human Capital Accumulation*.” *International Economic Review*, 45, 601–641.
- Indarte, Sasha (2023), “Moral Hazard versus Liquidity in Household Bankruptcy.” *Working Paper*.
- Jacob, Brian A, Damon Jones, and Benjamin J Keys (2023), “The Value of Student Debt Relief and the Role of Administrative Barriers: Evidence from the Teacher Loan Forgiveness Program.” *Working Paper*.

- Ji, Yan (2021), "Job Search under Debt: Aggregate Implications of Student Loans." *Journal of Monetary Economics*, 117, 741–759.
- JPMorgan Chase (2022), "Income Driven Repayment: Who needs student loan payment relief?" Technical report.
- Judd, Kenneth L. (1998), *Numerical Methods in Economics*. MIT Press, Cambridge, MA.
- Karamcheva, Nadia, Jeffrey Perry, and Constantine Yannelis (2020), "Income-Driven Repayment Plans for Student Loans." *Working Paper*.
- Kargar, Mahyar and William Mann (2022), "The Incidence of Student Loan Subsidies: Evidence from the PLUS Program." *The Review of Financial Studies*, hhac031.
- Karlan, Dean and Jonathan Zinman (2009), "Observing Unobservables: Identifying Information Asymmetries With a Consumer Credit Field Experiment." *Econometrica*, 77, 1993–2008.
- Keane, Michael and Richard Rogerson (2015), "Reconciling Micro and Macro Labor Supply Elasticities: A Structural Perspective." *Annual Review of Economics*, 7, 89–117.
- Keane, Michael P (2011), "Labor Supply and Taxes: A Survey." *Journal of Economic Literature*, 49, 961–1075.
- Keane, Michael P. (2016), "Life-cycle Labour Supply with Human Capital: Econometric and Behavioural Implications." *The Economic Journal*, 126, 546–577.
- Keane, Michael P. and Kenneth I. Wolpin (1997), "The Career Decisions of Young Men." *Journal of Political Economy*, 105, 473–523.
- Kleven, Henrik, Claus Kreiner, Kristian Larsen, and Jakob Søgaaard (2023), "Micro vs Macro Labor Supply Elasticities: The Role of Dynamic Returns to Effort." *Working Paper*.
- Kleven, Henrik J. and Mazhar Waseem (2013), "Using Notches to Uncover Optimization Frictions and Structural Elasticities: Theory and Evidence from Pakistan*." *The Quarterly Journal of Economics*, 128, 669–723.
- Lochner, Lance and A. Monge-Naranjo (2016), *Student Loans and Repayment: Theory, Evidence, and Policy*, volume 5. Elsevier B.V.
- Lucas, Robert E. (2003), "Macroeconomic Priorities." *American Economic Review*, 93, 1–14.
- Luo, Mi and Simon Mongey (2019), "Assets and Job Choice: Student Debt, Wages, and Amenities." *Working Paper*.
- Lusardi, Annamaria, Pierre Carl Michaud, and Olivia S. Mitchell (2017), "Optimal financial knowledge and wealth inequality." *Journal of Political Economy*, 125, 431–477.
- Makris, Miltiadis and Alessandro Pavan (2021), "Taxation under Learning by Doing." *Journal of Political Economy*, 129, 1878–1944.
- Marshall, Kate (2003), "Ease HECS burden on students, say universities." *Australian Financial Review*.
- Martin, Chelsey (2004), "For one in four, HECS now a lifelong debt." *Australian Financial Review*.
- Masatlioglu, Yusufcan and Efe A. Ok (2005), "Rational choice with status quo bias." *Journal of Economic Theory*, 121, 1–29.
- Matsuda, Kazushige and Karol Mazur (2022), "College education and income contingent loans in equilibrium." *Journal of Monetary Economics*, 132, 100–117.
- McFadden, Daniel (1989), "A Method of Simulated Moments for Estimation of Discrete Response Models Without Numerical Integration." *Econometrica*, 57, 995–1026.
- Medhora, Shalailah (2018), "From next year you'll be paying back your student loans sooner." *ABC News*.
- Mezza, Alvaro, Daniel Ringo, Shane Sherlund, and Kamila Sommer (2020), "Student loans and homeownership." *Journal of Labor Economics*, 38, 215–260.
- Mian, Atif and Amir Sufi (2014), *House of Debt: How They (and You) Caused the Great Recession, and How We Can Prevent It from Happening Again*. University of Chicago Press, Chicago, IL.
- Mincer, Jacob (1974), *Schooling, Experience, and Earnings*. National Bureau of Economic Research.
- Mirrlees, J. (1974), "Notes on Welfare Economics, Information and Uncertainty." In *Essays on Economic Behavior under Uncertainty*, Amsterdam: North Holland.

- Mueller, Holger M. and Constantine Yannelis (2019), “The rise in student loan defaults.” *Journal of Financial Economics*, 131, 1–19.
- Mueller, Holger M. and Constantine Yannelis (2021), “Reducing Barriers to Enrollment in Federal Student Loan Repayment Plans: Evidence from the Navient Field Experiment.” *Journal of Finance*.
- Mumford, Kevin J. (2022), “Student Selection into an Income Share Agreement.” *Working Paper*.
- Murto, Michael J (2022), “Student Loans and Human Capital Investments.” *Working Paper*.
- Nakamura, Emi and Jon Steinsson (2010), “Monetary Non-Neutrality in a Multi-Sector Menu Cost Model.” *Quarterly Journal of Economics*.
- Nelson, Brendan (2003), “Our universities: Backing Australia’s future.” Technical report, Commonwealth of Australia, Canberra, ACT.
- Nocedal, Jorge and Stephen J. Wright (2006), *Numerical Optimization*, 2nd ed edition. Springer Series in Operations Research, Springer, New York.
- Norton, Andrew (2018), “Has abolishing the discount for upfront payment of student contributions made a difference to upfront payment rates?” Technical report.
- Norton, Andrew (2019), “Demand-driven funding for universities is frozen. What does this mean and should the policy be restored?” *The Conversation*.
- Palacios, Miguel (2004), *Investing in Human Capital, A Capital Markets Approach to Student Funding*. Cambridge University Press.
- Piketty, Thomas and Emmanuel Saez (2013), “Optimal Labor Income Taxation.” In *Handbook of Public Economics*, volume 5, 391–474, Elsevier.
- Ramsey, F. P. (1927), “A Contribution to the Theory of Taxation.” *The Economic Journal*, 145, 47–61.
- Reuther, Albert, Jeremy Kepner, Chansup Byun, Siddharth Samsi, William Arcand, David Bestor, Bill Bergeron, Vijay Gadepally, Michael Houle, Matthew Hubbell, Michael Jones, Anna Klein, Lauren Milechin, Julia Mullen, Andrew Prout, Antonio Rosa, Charles Yee, and Peter Michaelias (2018), “Interactive supercomputing on 40,000 cores for machine learning and data analysis.” In *2018 IEEE High Performance Extreme Computing Conference (HPEC)*, 1–6.
- Robinson, Natasha (2019), “Documents reveal the Government looked at recovering HELP loans from deceased estates.” *ABC News*.
- Saez, Emmanuel (2001), “Using Elasticities to Derive Optimal Income Tax Rates.” *Review of Economic Studies*, 68, 205–229.
- Saez, Emmanuel (2002), “Optimal Income Transfer Programs: Intensive versus Extensive Labor Supply Responses.” *Quarterly Journal of Economics*, 1039–1074.
- Saez, Emmanuel (2010), “Do Taxpayers Bunch at Kink Points?” *American Economic Journal: Economic Policy*, 2, 180–212.
- Saez, Emmanuel, Joel Slemrod, and Seth H Giertz (2012), “The Elasticity of Taxable Income with Respect to Marginal Tax Rates: A Critical Review.” *Journal of Economic Literature*, 50, 3–50.
- Schmidt, Lawrence D. W. (2016), “Climbing and Falling Off the Ladder: Asset Pricing Implications of Labor Market Event Risk.” *Working Paper*.
- Shiller, Robert J. (2004), *The New Financial Order: Risks in the 21st Century*. Princeton University Press, Princeton, NJ.
- Slemrod, Joel (2019), “Tax Compliance and Enforcement.” *Journal of Economic Literature*, 57, 904–954.
- Slemrod, Joel and Shlomo Yitzhaki (2002), “Tax Avoidance, Evasion, and Administration.” *Handbook of Public Economics*, 3.
- Stantcheva, Stefanie (2017), “Optimal taxation and human capital policies over the life cycle.” *Journal of Political Economy*, 125, 1931–1990.
- Weil, Philippe (1990), “Unexpected Utility in Macroeconomics.” *Quarterly Journal of Economics*, 105, 29–42.
- Werquin, Nicolas (2015), “Income Taxation with Frictional Labor Supply.” *Working Paper*.
- Yannelis, Constantine (2020), “Strategic Default on Student Loans.” *Working Paper*.

Yannelis, Constantine and Greg Tracey (2022), “Student Loans and Borrower Outcomes.” *Annual Review of Financial Economics*.

Zeldes, Stephen P. (1989), “Optimal consumption with stochastic income: Deviations from certainty equivalence.” *Quarterly Journal of Economics*, 104, 275–297.

Zingales, Luigi (2012), “The College Graduate as Collateral.” *The New York Times*.

Required Disclaimer for Use of MADIP Data

The results of these studies are based, in part, on Australian Business Registrar (ABR) data supplied by the Registrar to the ABS under A New Tax System (Australian Business Number) Act 1999 and tax data supplied by the ATO to the ABS under the Taxation Administration Act 1953. These require that such data is only used for the purpose of carrying out functions of the ABS. No individual information collected under the Census and Statistics Act 1905 is provided back to the Registrar or ATO for administrative or regulatory purposes. Any discussion of data limitations or weaknesses is in the context of using the data for statistical purposes, and is not related to the ability of the data to support the ABR or ATO's core operational requirements. Legislative requirements to ensure privacy and secrecy of these data have been followed. Source data are de-identified and so data about specific individuals or firms has not been viewed in conducting this analysis. In accordance with the Census and Statistics Act 1905, results have been treated where necessary to ensure that they are not likely to enable identification of a particular person or organisation.